

# OPTIMIZING THE HYPERPARAMETERS OF ML (MLP, SVM, NAIVE BAYES) IN PREDICTING BREAST CANCER

## AIM

The project task is to analyze cell images and classify tumors as either cancerous or non-cancerous based on the features computed from the digitized images of a breast mass. The features are computed from the characteristics of the cell nuclei present in the image. The goal is to accurately determine if the tumor is malignant or benign.

Attribute Information(kaggle.com, n.d.):

id: ID number, diagnosis: The diagnosis of breast tissues (M = malignant, B = benign), radius\_mean: mean of distances from center to points on the perimeter, texture\_mean: standard deviation of gray-scale values, perimeter\_mean: mean size of the core tumor, area\_mean: area of the tumor, smoothness\_mean: mean of local variation in radius lengths, compactness\_mean: mean of  $\text{perimeter}^2 / \text{area} - 1.0$ , concavity\_mean: mean of severity of concave portions of the contour, concave\_points\_mean: mean for number of concave portions of the contour, symmetry\_mean, fractal\_dimension\_mean: mean for "coastline approximation" - 1, radius\_se: standard error for the mean of distances from center to points on the perimeter, texture\_se: standard error for standard deviation of gray-scale values, perimeter\_se, area\_se, smoothness\_se: standard error for local variation in radius lengths, compactness\_se: standard error for  $\text{perimeter}^2 / \text{area} - 1.0$ , concavity\_se: standard error for severity of concave portions of the contour, concave\_points\_se: standard error for number of concave portions of the contour, symmetry\_se, fractal\_dimension\_se: standard error for "coastline approximation" - 1, radius\_worst: "worst" or largest mean value for mean of distances from center to points on the perimeter, texture\_worst: "worst" or largest mean value for standard deviation of gray-scale values, perimeter\_worst, area\_worst, smoothness\_worst: "worst" or largest mean value for local variation in radius lengths, compactness\_worst: "worst" or largest mean value for  $\text{perimeter}^2 / \text{area} - 1.0$ , concavity\_worst: "worst" or largest mean value for severity of concave portions of the contour, concave\_points\_worst: "worst" or largest mean value for number of concave portions of the contour, symmetry\_worst, fractal\_dimension\_worst: "worst" or largest mean value for "coastline approximation" - 1

# Importing Dependencies

```
In [1]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout
```

2023-06-29 17:37:57.514551: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Data Collection & Loading

```
In [2]: df = pd.read_csv('/Users/oluwatoyineleja/Downloads/WISCONSIN.csv')
```

```
In [3]: #print the first 5 rows of dataframe
df.head()
```

Out[3]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	perimeter_min
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	101.0
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	101.0
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	101.0
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	101.0
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	101.0

5 rows × 33 columns

```
In [4]: #print last 5 rows of the dataframe
df.tail()
```

Out[4]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	perimeter_min
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	101.0
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	101.0
566	926954	M	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	101.0
567	927241	M	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	101.0
568	92751	B	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	101.0

5 rows × 33 columns

```
In [5]: #view number of rows and columns in the dataset
df.shape
```

```
Out[5]: (569, 33)
```

```
In [6]: #getting info about the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    569 non-null    int64
1   diagnosis                            569 non-null    object
2   radius_mean                          569 non-null    float64
3   texture_mean                         569 non-null    float64
4   perimeter_mean                       569 non-null    float64
5   area_mean                           569 non-null    float64
6   smoothness_mean                      569 non-null    float64
7   compactness_mean                     569 non-null    float64
8   concavity_mean                       569 non-null    float64
9   concave points_mean                  569 non-null    float64
10  symmetry_mean                        569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                            569 non-null    float64
14  perimeter_se                          569 non-null    float64
15  area_se                              569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
```

```

22 radius_worst      569 non-null    float64
23 texture_worst     569 non-null    float64
24 perimeter_worst   569 non-null    float64
25 area_worst        569 non-null    float64
26 smoothness_worst  569 non-null    float64
27 compactness_worst 569 non-null    float64
28 concavity_worst   569 non-null    float64
29 concave points_worst 569 non-null    float64
30 symmetry_worst    569 non-null    float64
31 fractal_dimension_worst 569 non-null    float64
32 Unnamed: 32       0 non-null      float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB

```

In [7]: *#statistical measures about the data*  
df.describe()

Out [7]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	poir
<b>count</b>	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569
<b>mean</b>	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	
<b>std</b>	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	
<b>min</b>	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	
<b>25%</b>	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	
<b>50%</b>	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	
<b>75%</b>	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	
<b>max</b>	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	

8 rows × 32 columns

## Preparing &amp; Cleaning the Data

```
In [8]: #checking for missing values  
df.isnull().sum()
```

```
Out[8]: id                0  
diagnosis                0  
radius_mean             0  
texture_mean            0  
perimeter_mean          0  
area_mean               0  
smoothness_mean         0  
compactness_mean        0  
concavity_mean           0  
concave points_mean     0  
symmetry_mean           0  
fractal_dimension_mean  0  
radius_se               0  
texture_se              0  
perimeter_se            0  
area_se                 0  
smoothness_se           0  
compactness_se          0  
concavity_se            0  
concave points_se       0  
symmetry_se             0  
fractal_dimension_se    0  
radius_worst            0  
texture_worst           0  
perimeter_worst         0  
area_worst              0  
smoothness_worst        0  
compactness_worst       0  
concavity_worst         0
```

```

concave points_worst      0
symmetry_worst            0
fractal_dimension_worst   0
Unnamed: 32               569
dtype: int64

```

Define Target and Input Variables

```

In [9]: #Rename Diagnosis to "Label" to easily identify our target variable
df = df.rename(columns = {'diagnosis':'label'})
print(df.dtypes)

```

```

id                int64
label             object
radius_mean       float64
texture_mean      float64
perimeter_mean    float64
area_mean         float64
smoothness_mean   float64
compactness_mean  float64
concavity_mean    float64
concave points_mean float64
symmetry_mean     float64
fractal_dimension_mean float64
radius_se         float64
texture_se        float64
perimeter_se      float64
area_se          float64
smoothness_se     float64
compactness_se    float64
concavity_se      float64
concave points_se float64
symmetry_se       float64
fractal_dimension_se float64

```

```
radius_worst      float64
texture_worst     float64
perimeter_worst   float64
area_worst        float64
smoothness_worst  float64
compactness_worst float64
concavity_worst   float64
concave points_worst float64
symmetry_worst    float64
fractal_dimension_worst float64
Unnamed: 32       float64
dtype: object
```

Separating the features and target variables

```
In [10]: #convert 'Target' categorical variable from [M]&[B] to [0]&[1] using Label Encoding
```

```
from sklearn.preprocessing import LabelEncoder

labelencoder = LabelEncoder()
Y = labelencoder.fit_transform(df["label"].values)
print("Label after encoding are: ", np.unique(Y))
```

```
Label after encoding are:  [0 1]
```

```
In [11]: #dropping column "id" and target "label"
X = df.drop(labels=["label", "id"], axis=1)
```

```
In [12]:
```



```
#view input variable
print (X)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	
...	...	...	...	...	...	
564	21.56	22.39	142.00	1479.0	0.11100	
565	20.13	28.25	131.20	1261.0	0.09780	
566	16.60	28.08	108.30	858.1	0.08455	
567	20.60	29.33	140.10	1265.0	0.11780	
568	7.76	24.54	47.92	181.0	0.05263	

	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	\
0	0.27760	0.30010	0.14710	0.2419	
1	0.07864	0.08690	0.07017	0.1812	
2	0.15990	0.19740	0.12790	0.2069	
3	0.28390	0.24140	0.10520	0.2597	
4	0.13280	0.19800	0.10430	0.1809	
...	...	...	...	...	
564	0.11590	0.24390	0.13890	0.1726	
565	0.10340	0.14400	0.09791	0.1752	
566	0.10230	0.09251	0.05302	0.1590	
567	0.27700	0.35140	0.15200	0.2397	
568	0.04362	0.00000	0.00000	0.1587	

	fractal_dimension_mean	...	texture_worst	perimeter_worst	area_worst	\
0	0.07871	...	17.33	184.60	2019.0	
1	0.05667	...	23.41	158.80	1956.0	
2	0.05999	...	25.53	152.50	1709.0	
3	0.09744	...	26.50	98.87	567.7	
4	0.05883	...	16.67	152.20	1575.0	

..	...	...	...	...
564	0.05623	...	26.40	166.10
565	0.05533	...	38.25	155.00
566	0.05648	...	34.12	126.70
567	0.07016	...	39.42	184.60
568	0.05884	...	30.37	59.16

	smoothness_worst	compactness_worst	concavity_worst	\
0	0.16220	0.66560	0.7119	
1	0.12380	0.18660	0.2416	
2	0.14440	0.42450	0.4504	
3	0.20980	0.86630	0.6869	
4	0.13740	0.20500	0.4000	
..	...	...	...	
564	0.14100	0.21130	0.4107	
565	0.11660	0.19220	0.3215	
566	0.11390	0.30940	0.3403	
567	0.16500	0.86810	0.9387	
568	0.08996	0.06444	0.0000	

	concave points_worst	symmetry_worst	fractal_dimension_worst	\
0	0.2654	0.4601	0.11890	
1	0.1860	0.2750	0.08902	
2	0.2430	0.3613	0.08758	
3	0.2575	0.6638	0.17300	
4	0.1625	0.2364	0.07678	
..	...	...	...	
564	0.2216	0.2060	0.07115	
565	0.1628	0.2572	0.06637	
566	0.1418	0.2218	0.07820	
567	0.2650	0.4087	0.12400	
568	0.0000	0.2871	0.07039	

Unnamed: 32

0	NaN
1	NaN

```
[569 rows x 31 columns]
```

[illegible]

```
In [14]: #CHECKING THE DISTRIBUTION OF TARGET VARIABLE  
df['label'].value_counts()
```


```
Out[14]: B    357  
        M    212  
        Name: label, dtype: int64
```

### Visualizing the Data

In this section we will develop some visualizations of the data in order to decide how to proceed with the multi-layer perceptron model and machine learning algorithms.

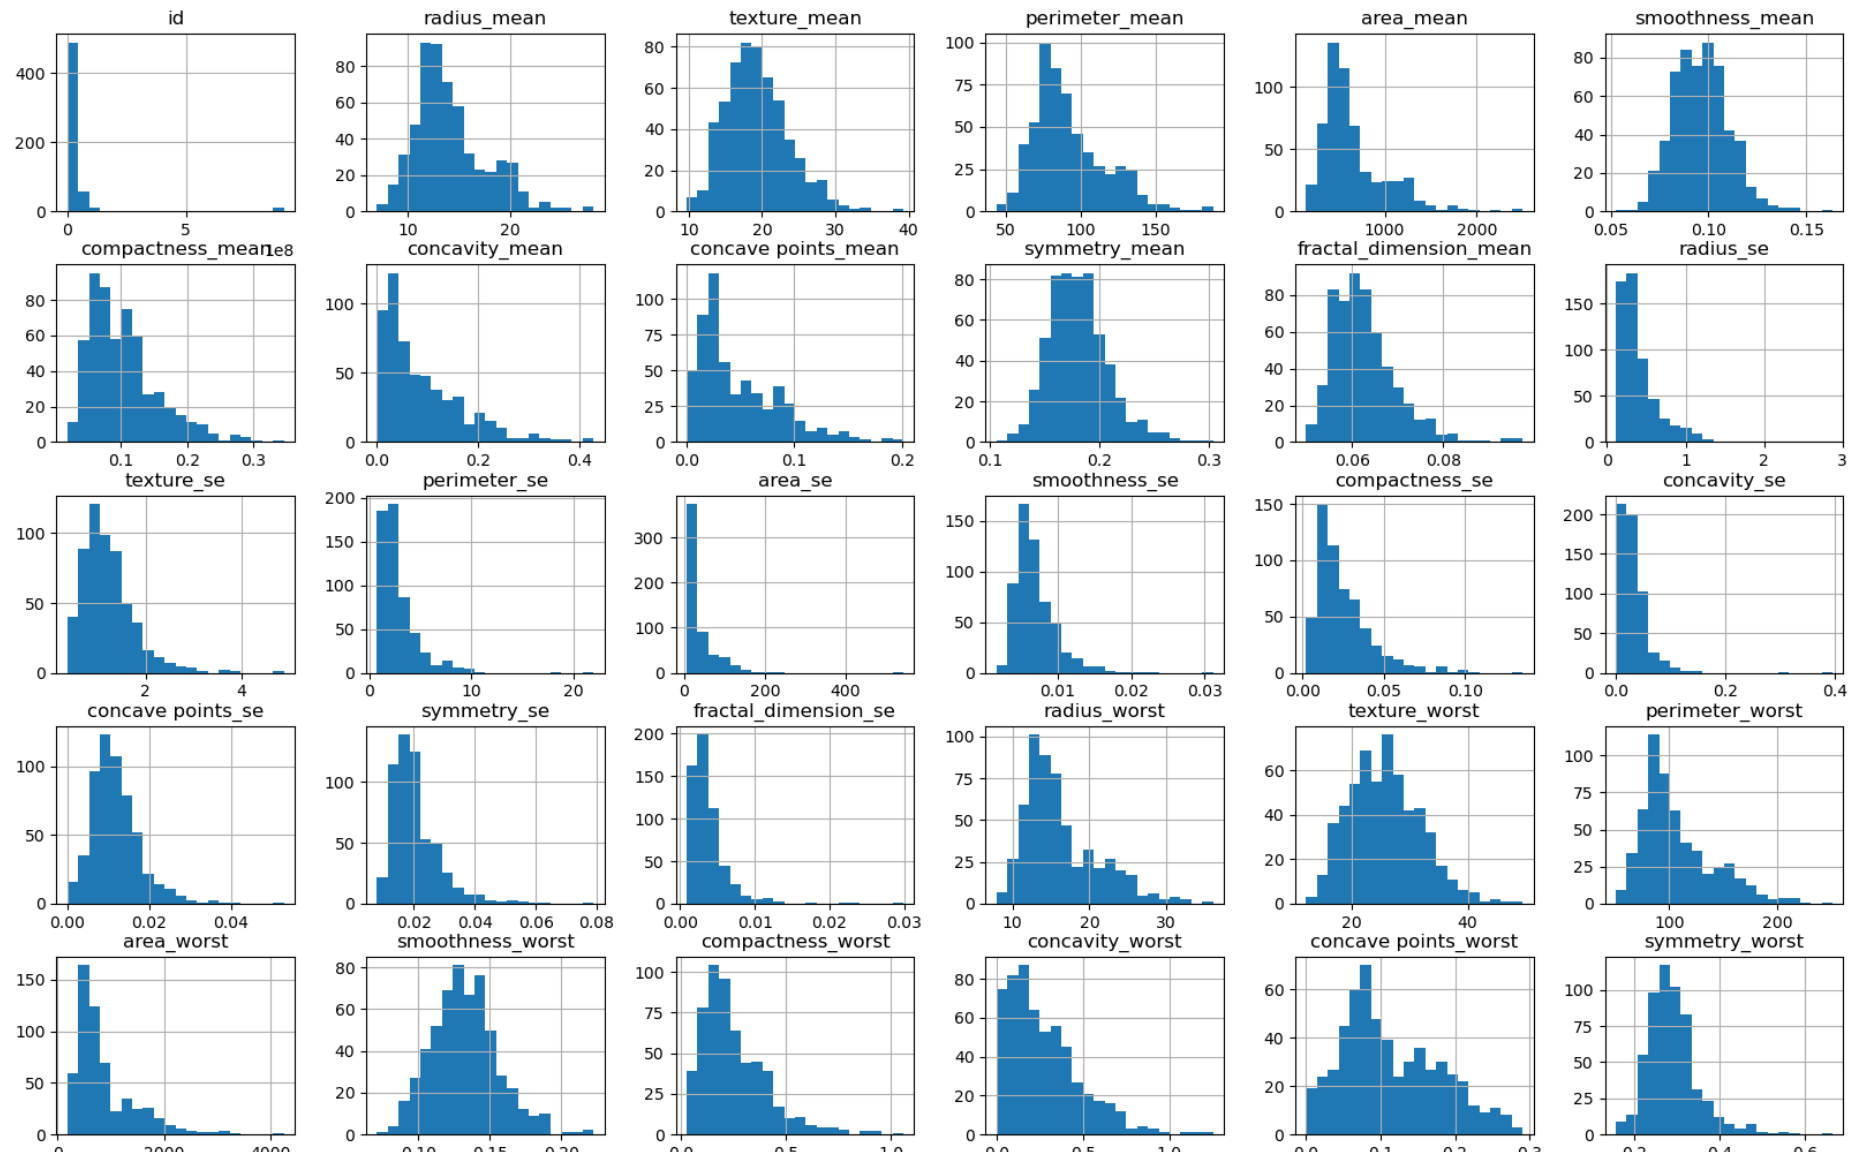
```
In [15]: #countplot to view how many benign and malignant cases present in dataset  
sns.countplot(data=df, x='label')
```

```
Out[15]: <AxesSubplot:xlabel='label', ylabel='count'>
```



```
In [16]:
```

```
#visualizing the distribution of each feature
df.hist(bins=20, figsize=(20,15))
plt.show()
```





Group input variables

```
In [17]: mean_features = list(df.columns[1:11])
se_features = list(df.columns[11:21])
worst_features = list(df.columns[21:31])
```

```
In [18]: #define pairplot's x and y axis
X_pairplot = X[["radius_mean", "texture_mean", "perimeter_mean", "area_mean", "smoothness_mean"]]
X_pairplot["label"] = Y

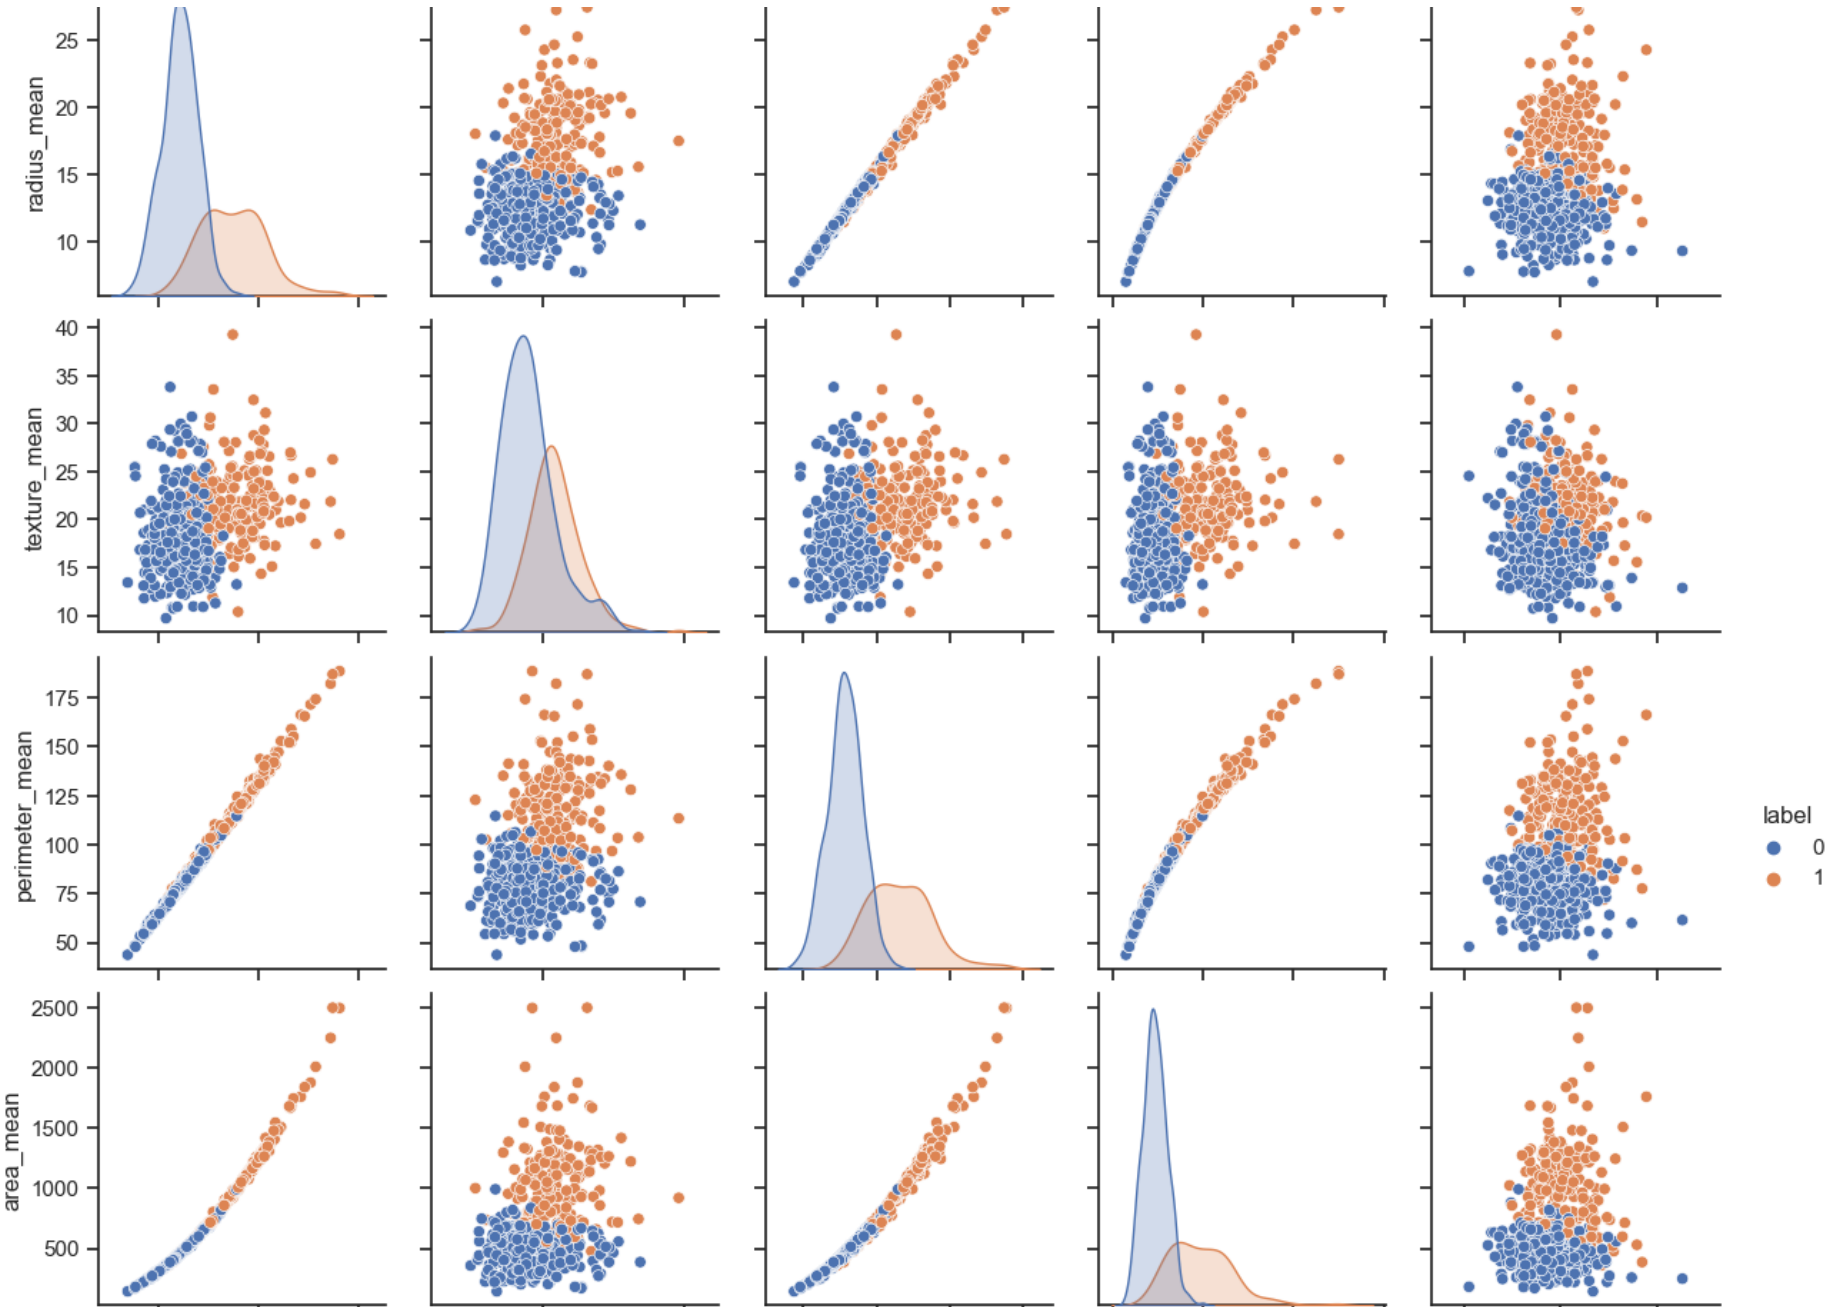
# Create a pairplot of the data
sns.set(style="ticks")
sns.pairplot(X_pairplot, hue="label", vars=["radius_mean", "texture_mean", "perimeter_mean", "area_mean", "smoothness_mean"])

/var/folders/vt/j0bzf6n54g9yf297psjl07w0000gn/T/ipykernel_7188/3861720049.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
X_pairplot["label"] = Y
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x7ff14b854f70>
```











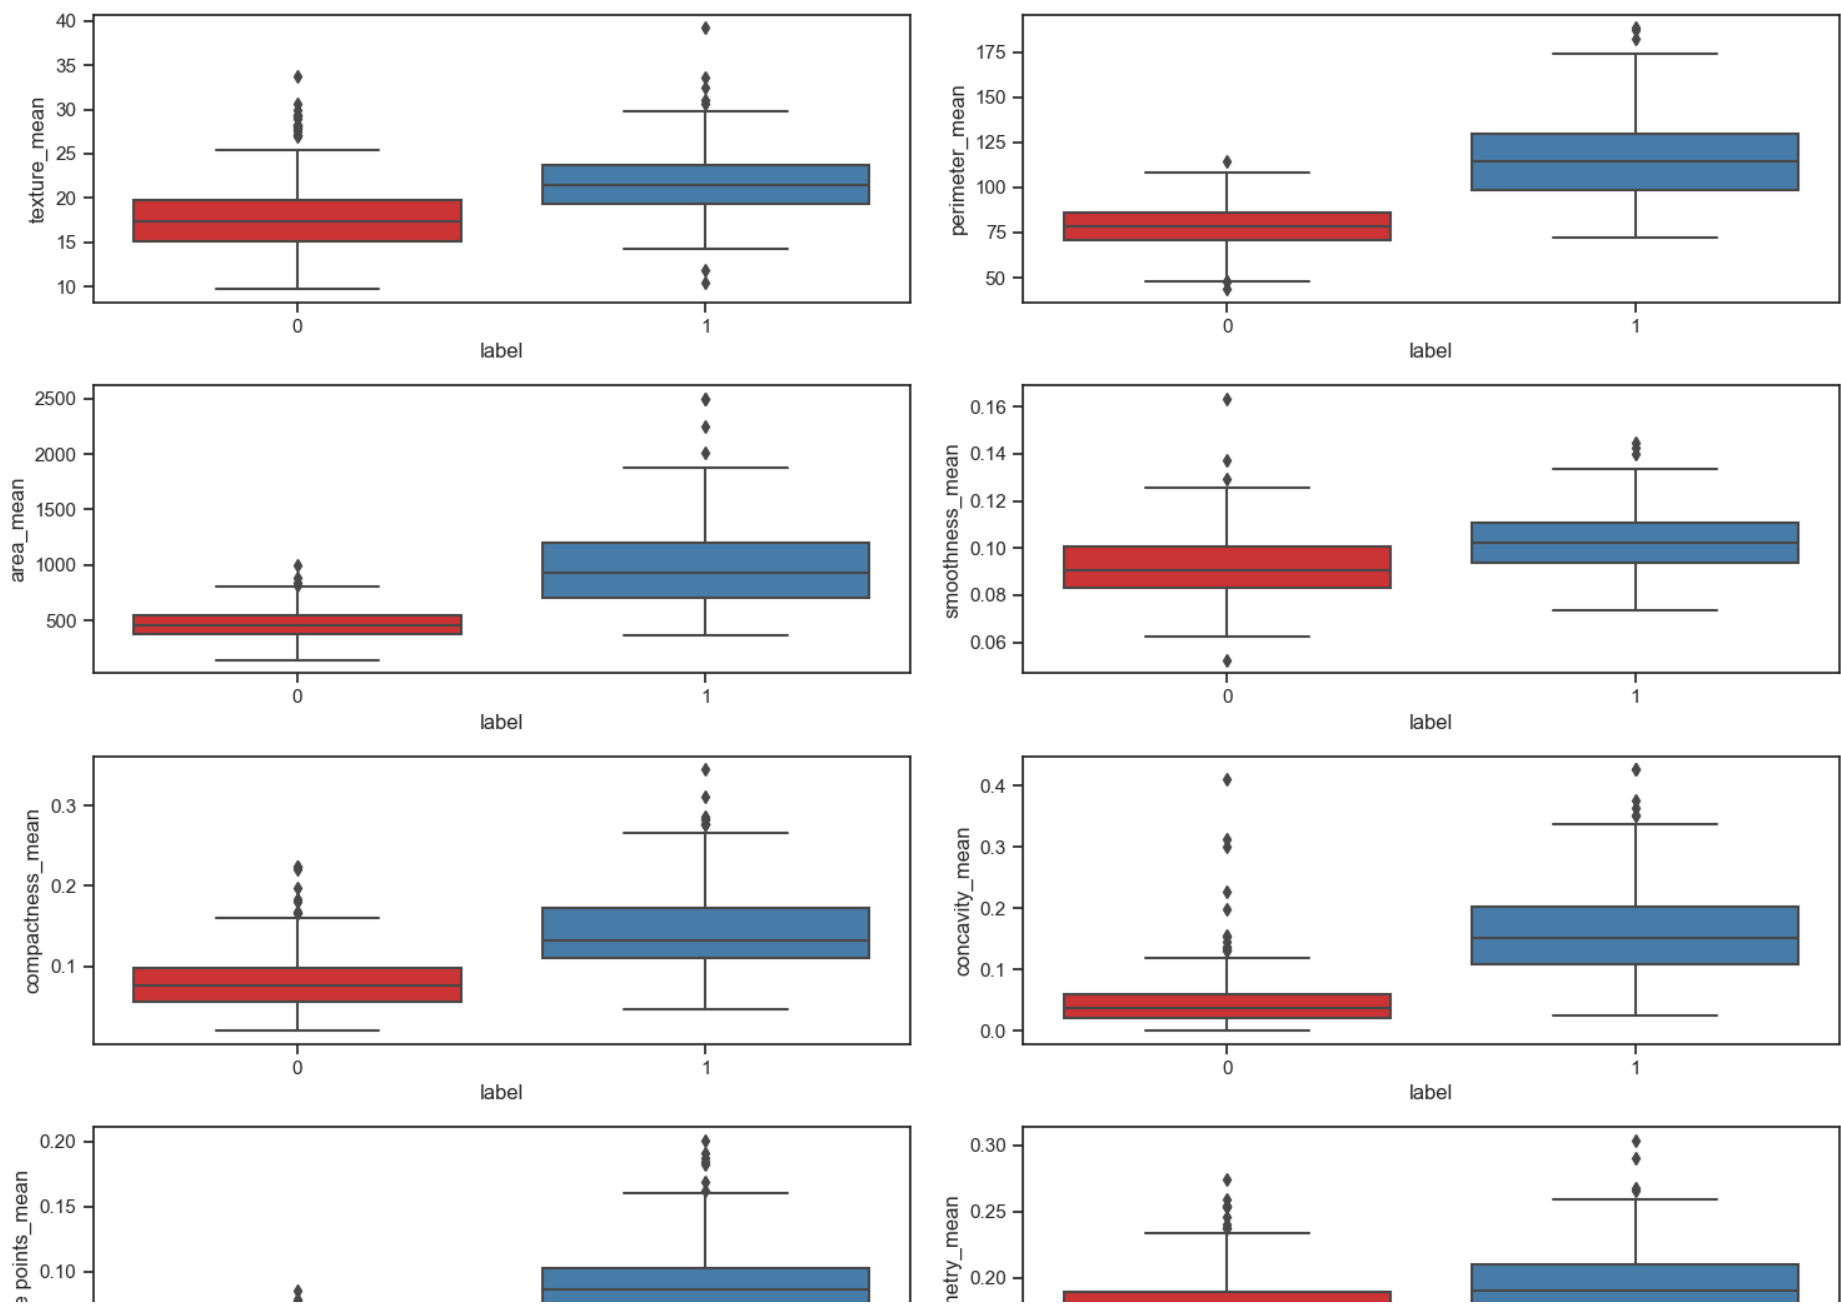
```
In [19]: # Combine the target variable and features into a single dataframe
data = pd.concat([pd.DataFrame(Y, columns=["label"]), X], axis=1)
plt.figure(figsize=(15,15))
for i, feature in enumerate(mean_features):
    rows = int(len(mean_features)/2)



    plt.subplot(rows, 2, 1+i)

    sns.boxplot(x='label', y=feature, data=data, palette='Set1')

plt.tight_layout()
plt.show()
```











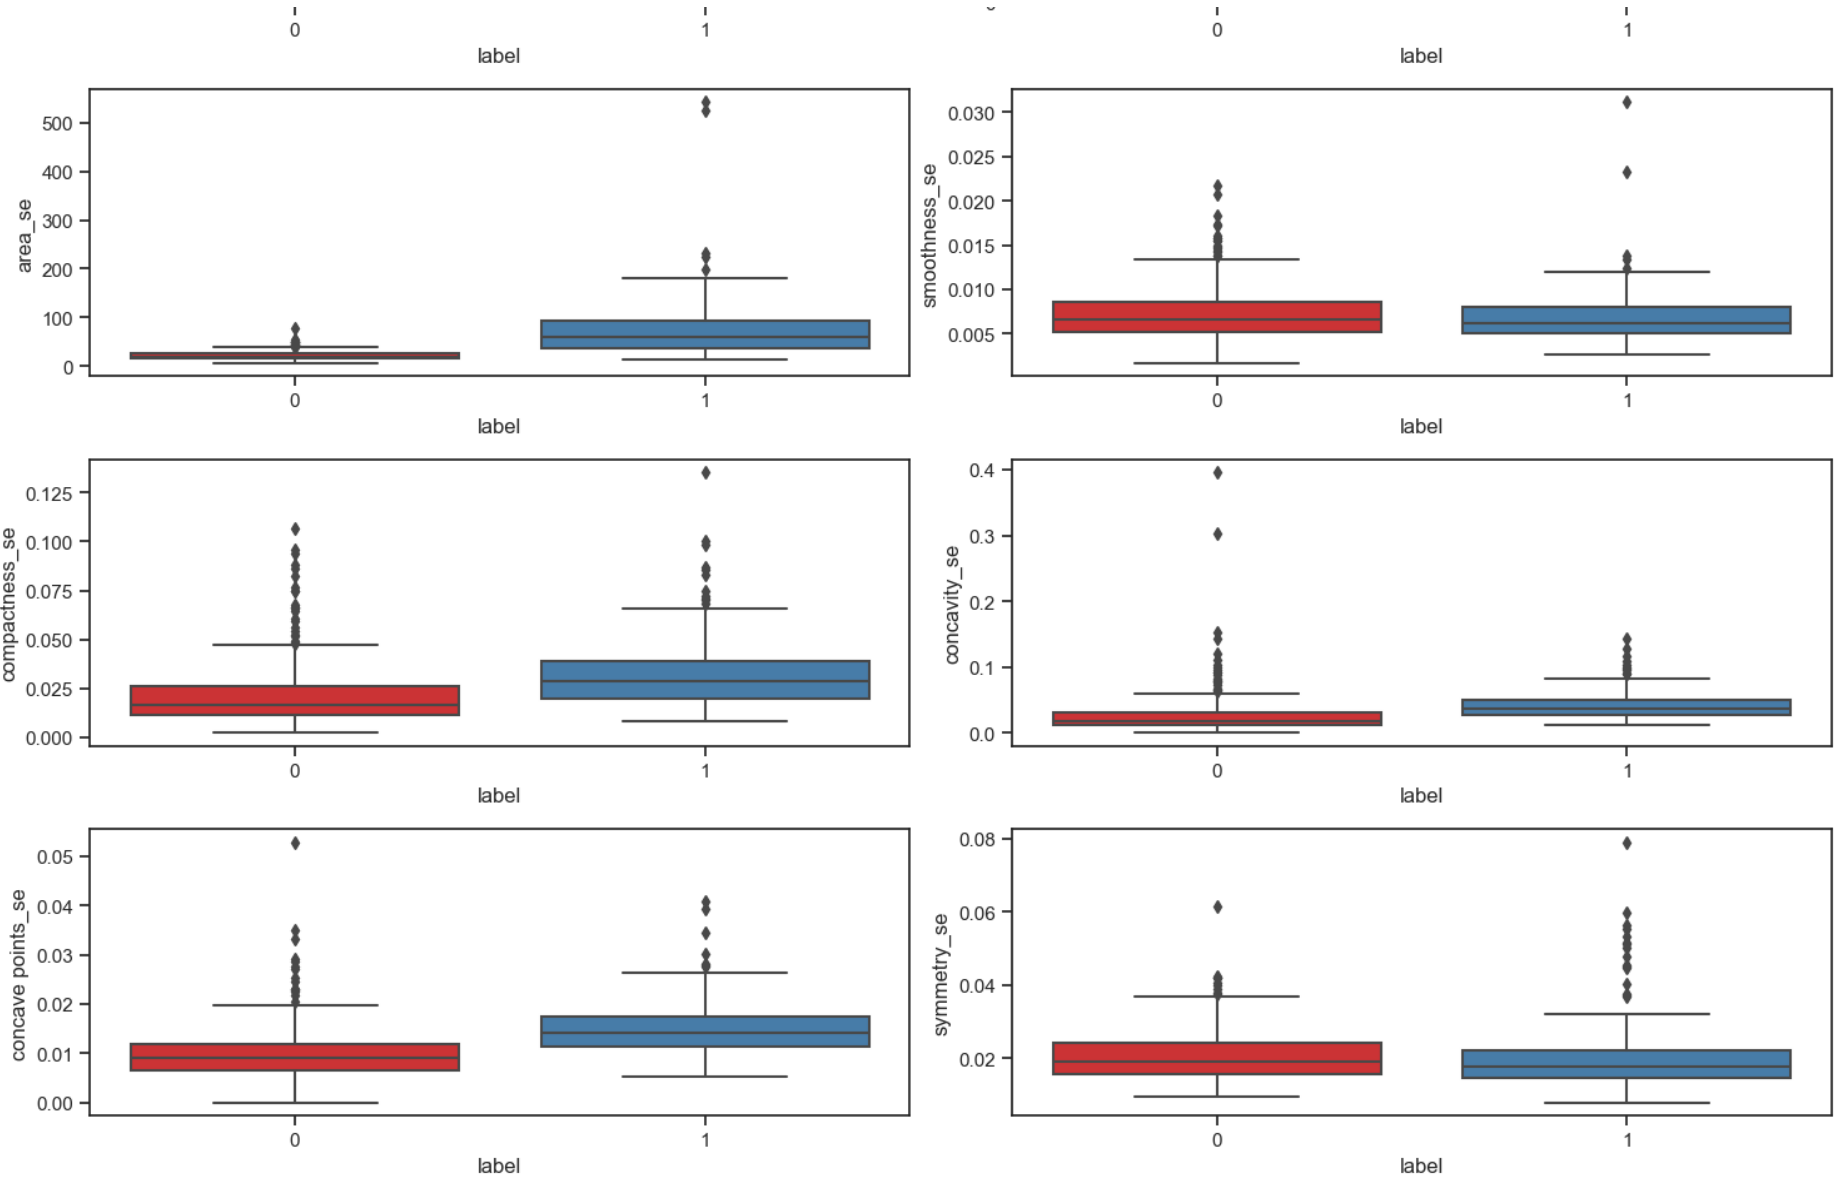
```
In [20]: plt.figure(figsize=(15,15))
for i, feature in enumerate(se_features):
    rows = int(len(mean_features)/2)

    plt.subplot(rows, 2, 1+i)

    sns.boxplot(x='label', y=feature, data=data, palette='Set1')

plt.tight_layout()
plt.show()
```










In [21]:

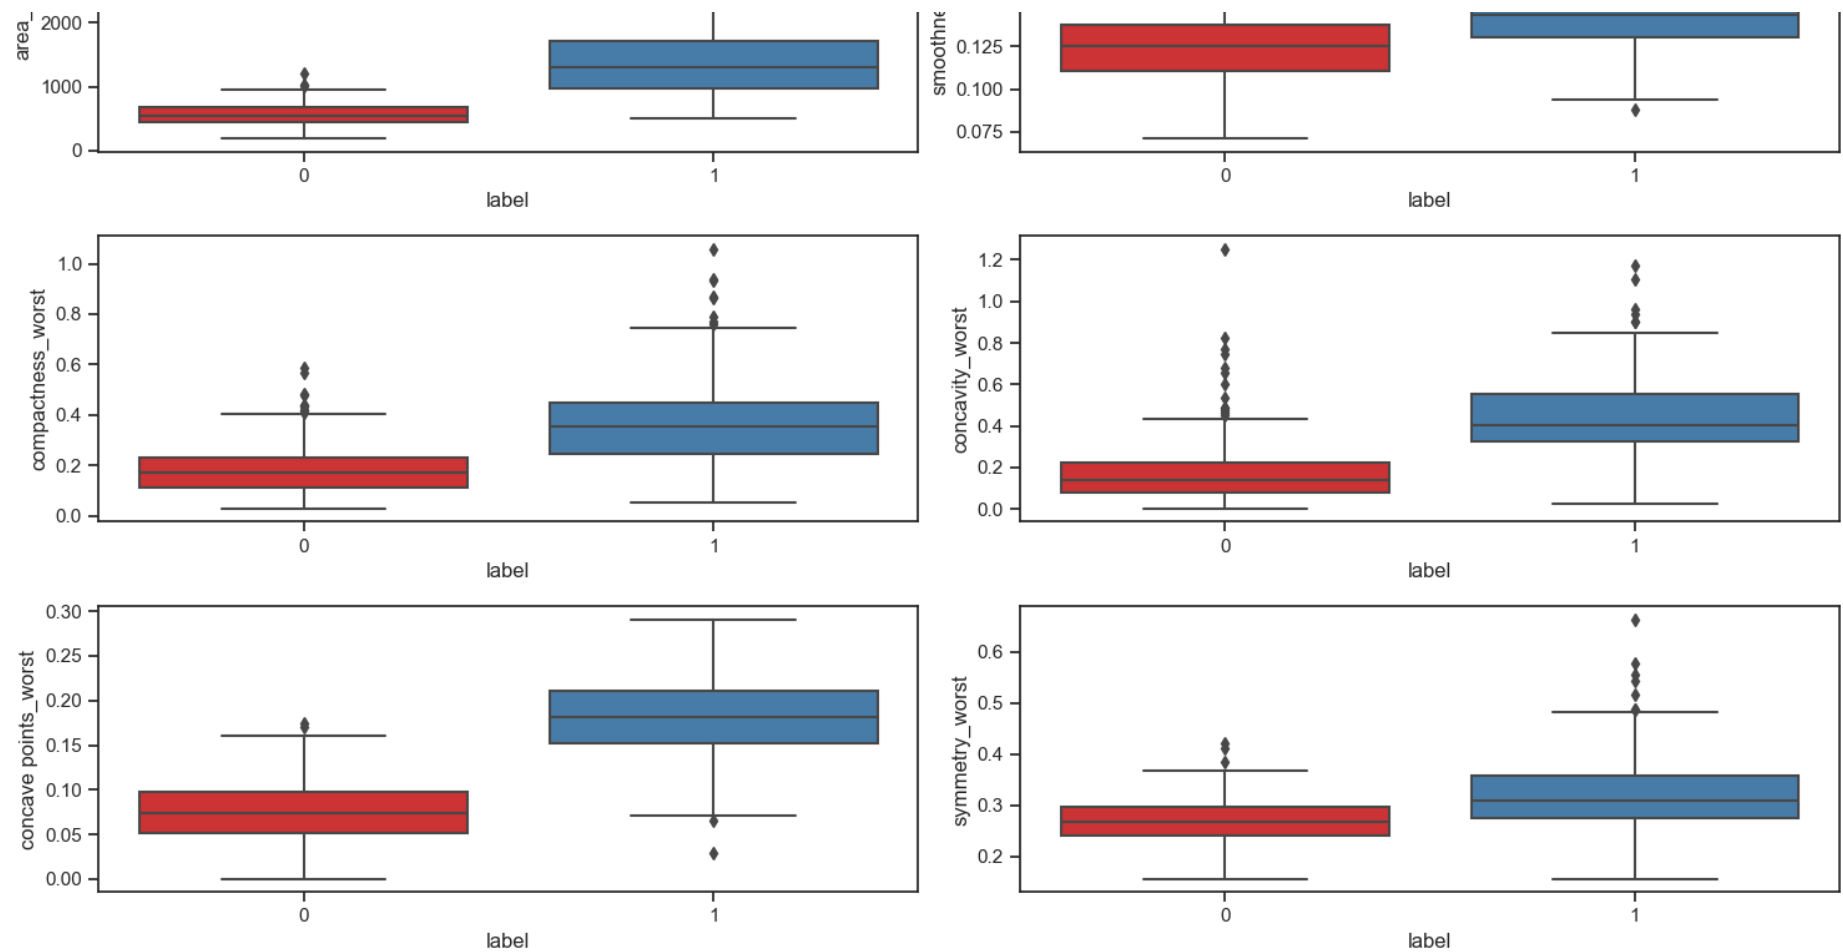
```
plt.figure(figsize=(15,15))
for i, feature in enumerate(worst_features):
    rows = int(len(mean_features)/2)

    plt.subplot(rows, 2, 1+i)

    sns.boxplot(x='label', y=feature, data=data, palette='Set1')

plt.tight_layout()
plt.show()
```





Explore Correlation

```
In [22]: #correlation of all features
data.corr()
```

Out[22]:

```
label    radius_mean    texture_mean    perimeter_mean    area_mean    smoothness_mean    compactness_mean    concavit
```

<b>label</b>	1.000000	0.730029	0.415185	0.742636	0.708984	0.358560	0.596534	C
<b>radius_mean</b>	0.730029	1.000000	0.323782	0.997855	0.987357	0.170581	0.506124	C
<b>texture_mean</b>	0.415185	0.323782	1.000000	0.329533	0.321086	-0.023389	0.236702	C
<b>perimeter_mean</b>	0.742636	0.997855	0.329533	1.000000	0.986507	0.207278	0.556936	C
<b>area_mean</b>	0.708984	0.987357	0.321086	0.986507	1.000000	0.177028	0.498502	C
<b>smoothness_mean</b>	0.358560	0.170581	-0.023389	0.207278	0.177028	1.000000	0.659123	C
<b>compactness_mean</b>	0.596534	0.506124	0.236702	0.556936	0.498502	0.659123	1.000000	C
<b>concavity_mean</b>	0.696360	0.676764	0.302418	0.716136	0.685983	0.521984	0.883121	1
<b>concave points_mean</b>	0.776614	0.822529	0.293464	0.850977	0.823269	0.553695	0.831135	C
<b>symmetry_mean</b>	0.330499	0.147741	0.071401	0.183027	0.151293	0.557775	0.602641	C
<b>fractal_dimension_mean</b>	-0.012838	-0.311631	-0.076437	-0.261477	-0.283110	0.584792	0.565369	C
<b>radius_se</b>	0.567134	0.679090	0.275869	0.691765	0.732562	0.301467	0.497473	C
<b>texture_se</b>	-0.008303	-0.097317	0.386358	-0.086761	-0.066280	0.068406	0.046205	C
<b>perimeter_se</b>	0.556141	0.674172	0.281673	0.693135	0.726628	0.296092	0.548905	C
<b>area_se</b>	0.548236	0.735864	0.259845	0.744983	0.800086	0.246552	0.455653	C
<b>smoothness_se</b>	-0.067016	-0.222600	0.006614	-0.202694	-0.166777	0.332375	0.135299	C
<b>compactness_se</b>	0.292999	0.206000	0.191975	0.250744	0.212583	0.318943	0.738722	C
<b>concavity_se</b>	0.253730	0.194204	0.143293	0.228082	0.207660	0.248396	0.570517	C
<b>concave points_se</b>	0.408042	0.376169	0.163851	0.407217	0.372320	0.380676	0.642262	C
<b>symmetry_se</b>	-0.006522	-0.104321	0.009127	-0.081629	-0.072497	0.200774	0.229977	C
<b>fractal_dimension_se</b>	0.077972	-0.042641	0.054458	-0.005523	-0.019887	0.283607	0.507318	C

<b>radius_worst</b>	0.776454	0.969539	0.352573	0.969476	0.962746	0.213120	0.535315	C
<b>texture_worst</b>	0.456903	0.297008	0.912045	0.303038	0.287489	0.036072	0.248133	C
<b>perimeter_worst</b>	0.782914	0.965137	0.358040	0.970387	0.959120	0.238853	0.590210	C
<b>area_worst</b>	0.733825	0.941082	0.343546	0.941550	0.959213	0.206718	0.509604	C
<b>smoothness_worst</b>	0.421465	0.119616	0.077503	0.150549	0.123523	0.805324	0.565541	C
<b>compactness_worst</b>	0.590998	0.413463	0.277830	0.455774	0.390410	0.472468	0.865809	C
<b>concavity_worst</b>	0.659610	0.526911	0.301025	0.563879	0.512606	0.434926	0.816275	C
<b>concave points_worst</b>	0.793566	0.744214	0.295316	0.771241	0.722017	0.503053	0.815573	C
<b>symmetry_worst</b>	0.416294	0.163953	0.105008	0.189115	0.143570	0.394309	0.510223	C
<b>fractal_dimension_worst</b>	0.323872	0.007066	0.119205	0.051019	0.003738	0.499316	0.687382	C
<b>Unnamed: 32</b>	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

32 rows × 32 columns


radius\_mean, perimeter\_mean, area\_mean have a high correlation with malignant tumor


```
In [23]: #correlation matrix of all features
plt.figure(figsize=(20,20))
sns.heatmap(data.corr(),annot=True)
```

Out[23]: <AxesSubplot:>










```
In [24]: #correlation matrix of mean features
plt.figure(figsize=(6,6))
sns.heatmap(df[mean_features].corr(),annot=True)
```

Out[24]: <AxesSubplot:>





Splitting data into training & test data.

```
In [25]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.25, random_state=42)
print("Shape of training data:", X_train.shape)
print("Shape of testing data:", X_test.shape)
```

Shape of training data: (426, 31)

Shape of testing data: (143, 31)

## FEATURES SELECTION (ON 11,20,30 FEATURES)

### 1) Removing highly correlated features & Randomly selecting features:

one or more important highly correlated features are removed to reduce multicollinearity. However, this could lead to loss of important information and hence, lower accuracy.

### 1) Training the model with low and randomly correlated 11 features

```
In [26]: Low_corr_vars = ['radius_mean', 'symmetry_mean', 'compactness_mean', 'texture_mean', 'compactness_worst',  
                          'symmetry_se', 'concavity_mean', 'texture_worst', 'concavity_se',  
                          'symmetry_worst', 'fractal_dimension_se']
```

```
In [27]: X_train = X_train[Low_corr_vars]  
X_test = X_test[Low_corr_vars]  
  
print("Shape of X_train:", X_train.shape)  
print("Shape of X_test:", X_test.shape)
```

Shape of X\_train: (426, 11)

Shape of X\_test: (143, 11)

Building the Neural Network

```
In [28]: #setting up layers of Neural network
model = Sequential()
model.add(Dense(16, input_dim=11, activation= 'relu'))
model.add(Dropout(0.2))
model.add(Dense(1))
model.add(Activation('sigmoid'))

#Compiling the neural network
model.compile(loss='binary_crossentropy', optimizer='adam', metrics = ['accuracy'])

print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 16)	192
dropout (Dropout)	(None, 16)	0
dense_1 (Dense)	(None, 1)	17
activation (Activation)	(None, 1)	0

=====  
Total params: 209

Trainable params: 209

Non-trainable params: 0

None

In [29]: *#training the neural network*

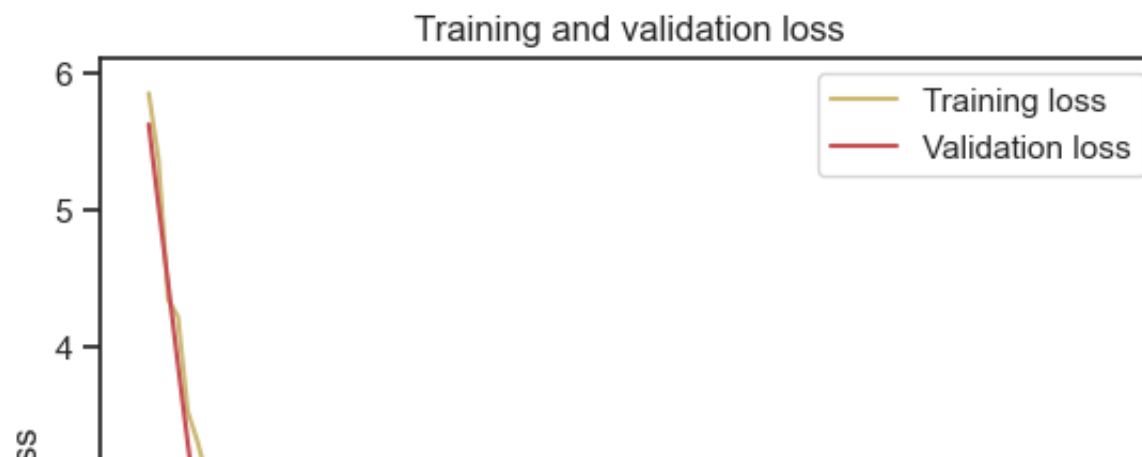
```
history = model.fit(X_train, y_train, verbose=1, epochs=100, batch_size=64, validation_data=(X_test, y_t
```


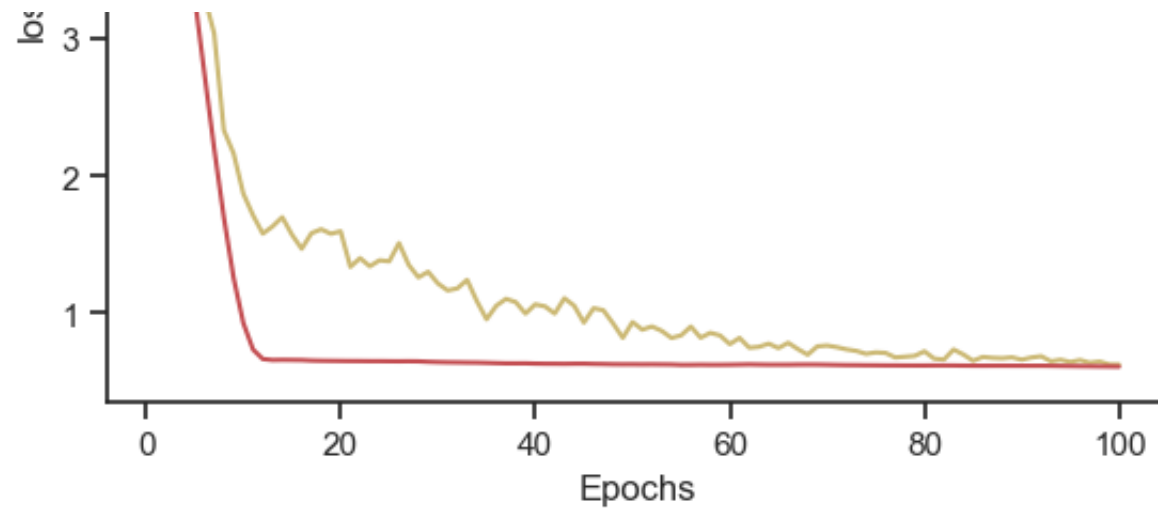
```
7/7 [=====] - 0s 10ms/step - loss: 1.3786 - accuracy: 0.5493 - val_loss: 0.6503 - val_accuracy: 0.6364
Epoch 26/100
7/7 [=====] - 0s 9ms/step - loss: 1.5111 - accuracy: 0.5376 - val_loss: 0.6498 - val_accuracy: 0.6364
Epoch 27/100
7/7 [=====] - 0s 9ms/step - loss: 1.3523 - accuracy: 0.5516 - val_loss: 0.6503 - val_accuracy: 0.6364
Epoch 28/100
7/7 [=====] - 0s 7ms/step - loss: 1.2617 - accuracy: 0.5657 - val_loss: 0.6495 - val_accuracy: 0.6364
Epoch 29/100
7/7 [=====] - 0s 8ms/step - loss: 1.3021 - accuracy: 0.5376 - val_loss: 0.6462 - val_accuracy: 0.6364
Epoch 30/100
7/7 [=====] - 0s 9ms/step - loss: 1.2171 - accuracy: 0.5516 - val_loss: 0.6442 - val_accuracy: 0.6294
Epoch 31/100
7/7 [=====] - 0s 9ms/step - loss: 1.1661 - accuracy: 0.5751 - val_loss: 0.6428 - val accuracy: 0.6364
```

In [30]:


```
#plot training and validation accuracy and loss at each epoch
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'y', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend()
plt.show()

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
plt.plot(epochs, acc, 'y', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```









```
In [31]: #predicting the test set results  
y_pred = model.predict(X_test)
```

```
y_pred = (y_pred > 0.5)
```


```
5/5 [=====] - 0s 2ms/step
```

```
In [32]: #confusion matrix

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True)
```

Out[32]: <AxesSubplot:>



## IMPACT OF LOW CORRELATED FEATURES

-The validation loss of a model is increased as these features might not be relevant to the target variable and can introduce noise to the model. -This can negatively impact the model's ability to generalize to new data, resulting in a higher validation loss. --Using low correlated features decreases the accuracy of the model, as these features may not be able to capture the important patterns and relationships in the data leading to the model performing poorly in predicting the target variable. -Using low correlated features led to an imbalanced confusion matrix, which skewed the model's classification towards one class with more false positives or false negatives. This can make it challenging to accurately classify the different classes and can result in a less balanced confusion matrix.

```
In [33]: from sklearn import metrics  
accuracy = metrics.accuracy_score(y_test, y_pred)
```

```
In [34]: #Evaluation of model
from sklearn import metrics
accuracy=metrics.accuracy_score(y_test,y_pred)
print ("The accuracy is %.2f" % accuracy)

#print the classification report
c_report=metrics.classification_report(y_test, y_pred)
print (c_report)
```


The accuracy is 0.69


	precision	recall	f1-score	support
0	0.67	0.97	0.79	89
1	0.80	0.22	0.35	54
accuracy			0.69	143
macro avg	0.74	0.59	0.57	143
weighted avg	0.72	0.69	0.62	143

Visualizing the Evaluation Result with low correlated values

In [35]:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
Accuracy = accuracy_score(y_test, y_pred)
# y_true is the true labels and y_pred is the corresponding predicted labels
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f1_score = 2 * precision * recall / (precision + recall)
Eval_Metrics = [Accuracy, precision, recall, f1_score]
Metric_Names = ['Accuracy', 'precision', 'recall', 'f1_score']
Metrics_pos = np.arange(len(Metric_Names))
plt.bar(Metrics_pos, Eval_Metrics)
plt.xticks(Metrics_pos, Metric_Names)
plt.title('Accuracy v Precision v Recall v F1 Score of the ANN model')
plt.show()
```





## IMPROVING THE MODEL

```
In [36]: #Standardizing the dataset
from sklearn.preprocessing import MinMaxScaler
```

2) Training the model with highly correlated 20 features & Standardizing the dataset

```
In [37]: High_corr_vars = ['radius_mean', 'perimeter_mean', 'area_mean', 'compactness_mean', 'concavity_mean',
                          'concave points_mean', 'radius_se', 'area_se', 'radius_worst', 'perimeter_worst',
                          'compactness_worst', 'smoothness_mean', 'symmetry_mean', 'perimeter_se', 'concavity_se',
                          'smoothness_worst', 'concavity_worst', 'area_worst', 'compactness_se', 'smoothness_mean']
x = df[High_corr_vars].values
```

```
In [38]: scaler = MinMaxScaler()
scaler.fit(x)
x = scaler.transform(x)
print (x)

[[0.52103744 0.54598853 0.36373277 ... 0.45069799 0.35139844 0.59375282]
 [0.64314449 0.61578329 0.50159067 ... 0.43521431 0.08132304 0.28987993]
 [0.60149557 0.59574321 0.44941676 ... 0.37450845 0.2839547 0.51430893]
 ...
 [0.45525108 0.44578813 0.30311771 ... 0.23073142 0.26330099 0.28816467]
 [0.64456434 0.66553797 0.4757158 ... 0.402035 0.44557936 0.58833619]
 [0.03686876 0.02853984 0.01590668 ... 0.02049744 0.01808514 0.          ]]
```

```
In [39]: #splitting to train and test set  
X_train, X_test, y_train, y_test = train_test_split(x, Y, test_size=0.25, random_state=42)
```

```
In [40]: print("Shape of X_train:", X_train.shape)  
print("Shape of X_test:", X_test.shape)
```

Shape of X\_train: (426, 20)

Shape of X\_test: (143, 20)

Building the neural network with: -Number of layers: 2

-Neurons per layer: Input layer: 20 Hidden layer: 16 Output layer: 1

-Activation function: Hidden layer: 'relu' Output layer: 'sigmoid'

-Dropout rate: 0.2 for the dropout layer

-Loss function: 'binary\_crossentropy'

-Optimizer: 'adam'

```
In [41]: #setting up layers of Neural network
model = Sequential()
model.add(Dense(16, input_dim=20, activation= 'relu'))
model.add(Dropout(0.2))
model.add(Dense(1))
model.add(Activation('sigmoid'))

#Compiling the neural network
model.compile(loss='binary_crossentropy', optimizer='adam', metrics = ['accuracy'])

print(model.summary())
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 16)	336
dropout_1 (Dropout)	(None, 16)	0
dense_3 (Dense)	(None, 1)	17
activation_1 (Activation)	(None, 1)	0
Total params: 353		
Trainable params: 353		
Non-trainable params: 0		
None		



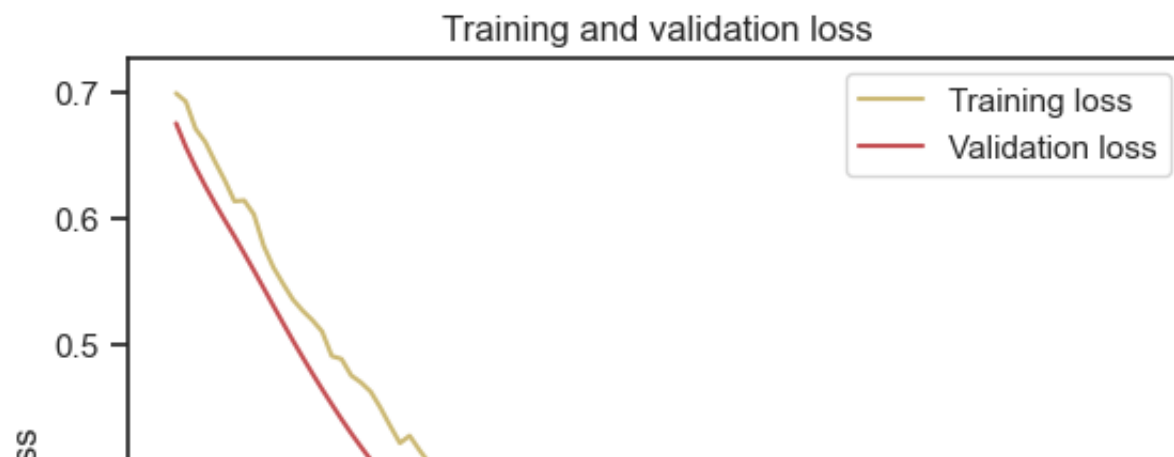
```
In [42]: history = model.fit(X_train, y_train, verbose=1, epochs=100, batch_size=64, validation_data=(X_test, y_test))
```



Epoch 29/100  
7/7 [=====] - 0s 10ms/step - loss: 0.3925 - accuracy: 0.8826 - val\_loss: 0.3376 - val\_accuracy: 0.9371  
Epoch 30/100  
7/7 [=====] - 0s 12ms/step - loss: 0.3764 - accuracy: 0.9108 - val\_loss: 0.3303 - val\_accuracy: 0.9371  
Epoch 31/100  
7/7 [=====] - 0s 10ms/step - loss: 0.3738 - accuracy: 0.9085 - val\_loss: 0.3229 - val\_accuracy: 0.9371  
Epoch 32/100  
7/7 [=====] - 0s 15ms/step - loss: 0.3559 - accuracy: 0.9131 - val\_loss: 0.3158 - val\_accuracy: 0.9371  
Epoch 33/100  
7/7 [=====] - 0s 13ms/step - loss: 0.3535 - accuracy: 0.9061 - val\_loss: 0.3090 - val\_accuracy: 0.9371  
Epoch 34/100  
7/7 [=====] - 0s 11ms/step - loss: 0.3518 - accuracy: 0.9014 - val\_loss: 0.3024 - val\_accuracy: 0.9371  
Epoch 35/100  
7/7 [=====] - 0s 13ms/step - loss: 0.3300 - accuracy: 0.9155 - val\_loss: 0.2906 - val\_accuracy: 0.9371


```
In [43]:
```

```
#plot training and validation accuracy and loss at each epoch
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'y', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('loss')
plt.legend()
plt.show()

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
plt.plot(epochs, acc, 'y', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```







```
In [44]: #predicting the test set results
y_pred = model.predict(X_test)
```

```
y_pred = (y_pred > 0.5)
```


```
5/5 [=====] - 0s 2ms/step
```

```
In [45]: #confusion matrix

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True)
```

Out[45]: <AxesSubplot:>



```
In [46]: accuracy = accuracy_score(y_test, y_pred)
print("Testing accuracy:", accuracy)
```

Testing accuracy: 0.972027972027972

```
In [47]: training_accuracy = history.history['accuracy'][-1]
print("Training accuracy:", training_accuracy)
```

Training accuracy: 0.9413145780563354

The testing data accuracy is higher than the accuracy on the training data, it suggests that the model is good at making predictions on new and unseen data, and not just the data that it was trained on. However, having a high testing accuracy is not a guarantee that the model is perfect or will work well in all situations. It's possible that the model may have been overfitted to the testing data or may perform poorly on data that it has not seen before. To ensure the model's reliability, I evaluated its performance using different evaluation metrics, rather than making any hasty conclusions about its performance.

```
In [48]: #Evaluation of model
from sklearn import metrics
accuracy=metrics.accuracy_score(y_test,y_pred)
print ("The accuracy is %.2f" % accuracy)


#print the classification report
c_report=metrics.classification_report(y_test, y_pred)
print (c_report)
```

The accuracy is 0.97


	precision	recall	f1-score	support
0	0.97	0.99	0.98	89
1	0.98	0.94	0.96	54
accuracy			0.97	143
macro avg	0.97	0.97	0.97	143
weighted avg	0.97	0.97	0.97	143

In [49]:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
Accuracy = accuracy_score(y_test, y_pred)
# y_true is the true labels and y_pred is the corresponding predicted labels
tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()
precision = tp / (tp + fp)
recall = tp / (tp + fn)
f1_score = 2 * precision * recall / (precision + recall)
Eval_Metrics = [Accuracy, precision, recall, f1_score]
Metric_Names = ['Accuracy', 'precision', 'recall', 'f1_score']
Metrics_pos = np.arange(len(Metric_Names))
plt.bar(Metrics_pos, Eval_Metrics)
plt.xticks(Metrics_pos, Metric_Names)
plt.title('Accuracy v Precision v Recall v F1 Score of the ANN model')
plt.show()
```







### 3) TRAINING THE MODEL WITH 30 FEATURES

```
In [50]: #Define input variable as P
P = df.drop(labels=["label", "id"], axis=1)
print(P.describe().T)
```

	count	mean	std	min	\
radius_mean	569.0	14.127292	3.524049	6.981000	
texture_mean	569.0	19.289649	4.301036	9.710000	
perimeter_mean	569.0	91.969033	24.298981	43.790000	
area_mean	569.0	654.889104	351.914129	143.500000	
smoothness_mean	569.0	0.096360	0.014064	0.052630	
compactness_mean	569.0	0.104341	0.052813	0.019380	
concavity_mean	569.0	0.088799	0.079720	0.000000	
concave points_mean	569.0	0.048919	0.038803	0.000000	
symmetry_mean	569.0	0.181162	0.027414	0.106000	
fractal_dimension_mean	569.0	0.062798	0.007060	0.049960	
radius_se	569.0	0.405172	0.277313	0.111500	
texture_se	569.0	1.216853	0.551648	0.360200	
perimeter_se	569.0	2.866059	2.021855	0.757000	
area_se	569.0	40.337079	45.491006	6.802000	
smoothness_se	569.0	0.007041	0.003003	0.001713	
compactness_se	569.0	0.025478	0.017908	0.002252	
concavity_se	569.0	0.031894	0.030186	0.000000	
concave points_se	569.0	0.011796	0.006170	0.000000	
symmetry_se	569.0	0.020542	0.008266	0.007882	
fractal_dimension_se	569.0	0.003795	0.002646	0.000895	
radius_worst	569.0	16.269190	4.833242	7.930000	

texture_worst	569.0	25.677223	6.146258	12.020000
perimeter_worst	569.0	107.261213	33.602542	50.410000
area_worst	569.0	880.583128	569.356993	185.200000
smoothness_worst	569.0	0.132369	0.022832	0.071170
compactness_worst	569.0	0.254265	0.157336	0.027290
concavity_worst	569.0	0.272188	0.208624	0.000000
concave points_worst	569.0	0.114606	0.065732	0.000000
symmetry_worst	569.0	0.290076	0.061867	0.156500
fractal_dimension_worst	569.0	0.083946	0.018061	0.055040
Unnamed: 32	0.0	NaN	NaN	NaN

	25%	50%	75%	max
radius_mean	11.700000	13.370000	15.780000	28.11000
texture_mean	16.170000	18.840000	21.800000	39.28000
perimeter_mean	75.170000	86.240000	104.100000	188.50000
area_mean	420.300000	551.100000	782.700000	2501.00000
smoothness_mean	0.086370	0.095870	0.105300	0.16340
compactness_mean	0.064920	0.092630	0.130400	0.34540
concavity_mean	0.029560	0.061540	0.130700	0.42680
concave points_mean	0.020310	0.033500	0.074000	0.20120
symmetry_mean	0.161900	0.179200	0.195700	0.30400
fractal_dimension_mean	0.057700	0.061540	0.066120	0.09744
radius_se	0.232400	0.324200	0.478900	2.87300
texture_se	0.833900	1.108000	1.474000	4.88500
perimeter_se	1.606000	2.287000	3.357000	21.98000
area_se	17.850000	24.530000	45.190000	542.20000
smoothness_se	0.005169	0.006380	0.008146	0.03113
compactness_se	0.013080	0.020450	0.032450	0.13540
concavity_se	0.015090	0.025890	0.042050	0.39600
concave points_se	0.007638	0.010930	0.014710	0.05279
symmetry_se	0.015160	0.018730	0.023480	0.07895
fractal_dimension_se	0.002248	0.003187	0.004558	0.02984
radius_worst	13.010000	14.970000	18.790000	36.04000
texture_worst	21.080000	25.410000	29.720000	49.54000
perimeter_worst	84.110000	97.660000	125.400000	251.20000

area_worst	515.300000	686.500000	1084.000000	4254.000000
smoothness_worst	0.116600	0.131300	0.146000	0.22260
compactness_worst	0.147200	0.211900	0.339100	1.05800
concavity_worst	0.114500	0.226700	0.382900	1.25200
concave points_worst	0.064930	0.099930	0.161400	0.29100
symmetry_worst	0.250400	0.282200	0.317900	0.66380
fractal_dimension_worst	0.071460	0.080040	0.092080	0.20750
Unnamed: 32	NaN	NaN	NaN	NaN

In [51]: *#standardize input data*

```
scaled = MinMaxScaler()
scaled.fit(P)
P = scaled.transform(P)
print (P)
```

```
[[0.52103744 0.0226581 0.54598853 ... 0.59846245 0.41886396 nan]
 [0.64314449 0.27257355 0.61578329 ... 0.23358959 0.22287813 nan]
 [0.60149557 0.3902604 0.59574321 ... 0.40370589 0.21343303 nan]
 ...
 [0.45525108 0.62123774 0.44578813 ... 0.12872068 0.1519087 nan]
 [0.64456434 0.66351031 0.66553797 ... 0.49714173 0.45231536 nan]
 [0.03686876 0.50152181 0.02853984 ... 0.25744136 0.10068215 nan]]
```

/opt/anaconda3/lib/python3.9/site-packages/sklearn/preprocessing/\_data.py:461: RuntimeWarning: All-NaN slice encountered

```
data_min = np.nanmin(X, axis=0)
```

/opt/anaconda3/lib/python3.9/site-packages/sklearn/preprocessing/\_data.py:462: RuntimeWarning: All-NaN slice encountered

```
data_max = np.nanmax(X, axis=0)
```

```
In [52]: #split training and test data  
X_train, X_test, y_train, y_test = train_test_split(P, Y, test_size=0.25, random_state=42)  
print("Shape of training data:", X_train.shape)  
print("Shape of testing data:", X_test.shape)
```

```
Shape of training data: (426, 31)  
Shape of testing data: (143, 31)
```

```
In [53]: #checking for the expected size of model  
print("Expected input shape of model:", model.layers[0].input_shape)
```

```
Expected input shape of model: (None, 20)
```

```
In [54]: #remove a variable to adjust input data size to match model size  
X_train = X_train[:, :-1]  
X_test = X_test[:, :-1]
```

```
In [55]: print("Shape of X_train:", X_train.shape)  
print("Shape of X_test:", X_test.shape)
```

```
Shape of X_train: (426, 30)  
Shape of X_test: (143, 30)
```

Here are the hyperparameters used in the below code:

Number of layers: 3

Neurons per layer: Input layer: 30 Hidden layers: 32, 16 Output layer: 1

Activation function: Hidden layers: 'relu' Output layer: 'sigmoid'

Dropout rate: 0.5 for each dropout layer

Loss function: 'binary\_crossentropy'

Optimizer: 'rmsprop'

Evaluation metric: 'accuracy'


verbose: 1

batch\_size: 32

validation\_split: 0.2

Epoch:50

In [56]:



```

#Build neural network model

import keras.backend as K
K.clear_session()

#setting up layers of Neural network
model = Sequential()
model.add(Dense(32, input_dim=30, activation= 'relu'))
model.add(Dropout(0.5))
model.add(Dense(16, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

#Compiling the neural network
model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics = ['accuracy'])

print(model.summary())

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 32)	992
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 16)	528
dropout_1 (Dropout)	(None, 16)	0
dense_2 (Dense)	(None, 1)	17
=====		
Total params: 1,537		
Trainable params: 1,537		

Non-trainable params: 0

None

```
In [57]: history = model.fit(X_train, y_train, verbose=1, epochs=50, batch_size=32, validation_split=0.2)
245 - val_accuracy: 0.8953
Epoch 11/50
11/11 [=====] - 0s 8ms/step - loss: 0.5459 - accuracy: 0.7882 - val_loss: 0.50
27 - val_accuracy: 0.8953
Epoch 12/50
11/11 [=====] - 0s 7ms/step - loss: 0.5400 - accuracy: 0.7941 - val_loss: 0.48
28 - val_accuracy: 0.9070
Epoch 13/50
11/11 [=====] - 0s 8ms/step - loss: 0.5305 - accuracy: 0.7912 - val_loss: 0.46
28 - val_accuracy: 0.9070
Epoch 14/50
11/11 [=====] - 0s 7ms/step - loss: 0.5374 - accuracy: 0.7735 - val_loss: 0.44
66 - val_accuracy: 0.9070
Epoch 15/50
11/11 [=====] - 0s 8ms/step - loss: 0.5061 - accuracy: 0.8118 - val_loss: 0.43
11 - val_accuracy: 0.8953
Epoch 16/50
11/11 [=====] - 0s 8ms/step - loss: 0.4620 - accuracy: 0.8118 - val_loss: 0.40
94 - val_accuracy: 0.9070
Epoch 17/50
```

```
In [58]: #predicting the test set results
y_pred = model.predict(X_test)

y_pred = (y_pred > 0.5)


5/5 [=====] - 0s 2ms/step
```

```
In [59]: #confusion matrix

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True)
```

Out[59]: <AxesSubplot:>





```
In [60]: accuracy = accuracy_score(y_test, y_pred)
print("Testing accuracy:", accuracy)
```

Testing accuracy: 0.958041958041958

```
In [61]: training_accuracy = history.history['accuracy'][-1]
print("Training accuracy:", training_accuracy)
```

Training accuracy: 0.9413145780563354

```
In [62]: #Evaluation of model
from sklearn import metrics
accuracy=metrics.accuracy_score(y_test,y_pred)
print ("The accuracy is %.2f" % accuracy)

#print the classification report
c_report=metrics.classification_report(y_test, y_pred)
print (c_report)
```

The accuracy is 0.96

	precision	recall	f1-score	support
0	0.96	0.98	0.97	89
1	0.96	0.93	0.94	54
accuracy			0.96	143
macro avg	0.96	0.95	0.96	143
weighted avg	0.96	0.96	0.96	143

```
In [63]: # Generating data for input
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)
```

Shape of X\_train: (426, 30)

Shape of X\_test: (143, 30)

## HYPERPARAMETERS TUNING

This consist of various parameters listed below

- 1.Number of layers: The number of layers in the neural network architecture.
- 2.Neurons per layer: The number of neurons in each layer of the neural network architecture.
- 3.Activation function: The activation function is used to introduce non-linearity to the output of each neuron in the neural network.
- 4.Dropout rate: Dropout is a regularization technique used to reduce overfitting in the neural network. Dropout layers randomly drop out a fraction of the neurons during training.
- 5.Loss function: The loss function is used to measure how well the neural network is performing during training. In this case, the binary cross-entropy loss function has been used as the data is binary classification.
- 6.Optimizer: The optimizer is the algorithm used to update the weights and biases of the neural network during training in order to minimize the loss function.
- 7.Verbose: 0, 1 or 2. Verbosity mode, 0 = silent, 1 = progress bar, 2 = one line per epoch.
- 8.Epochs: the number of times to iterate over the entire training dataset
- 9.Batch\_size: the number of samples per batch of training data
- 10.Validation\_split: fraction of the training data to use as validation data. For example, if set to 0.2, 20% of the training data will be used as validation data.

## VARIOUS HYPERPARAMETER ACHICTECTURE

In [64]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.metrics import classification_report
from tensorflow.keras import regularizers
import tensorflow as tf

# Set the random seed for reproducibility
seed_value = 42
np.random.seed(seed_value)
tf.random.set_seed(seed_value)

# Architecture 1: 1 hidden layer with 4 neurons
model1 = Sequential()
model1.add(Dense(4, input_dim=30, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model1.add(Dropout(0.1))
model1.add(Dense(1, activation='sigmoid'))

# Architecture 2: 2 hidden layers with 4 and 8 neurons
model2 = Sequential()
model2.add(Dense(4, input_dim=30, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model2.add(Dropout(0.2))
model2.add(Dense(8, activation='relu'))
model2.add(Dense(1, activation='sigmoid'))

# Architecture 3: 2 hidden layers with 8 neurons each
model3 = Sequential()
model3.add(Dense(8, input_dim=30, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model3.add(Dropout(0.3))
model3.add(Dense(8, activation='relu'))
model3.add(Dense(1, activation='sigmoid'))

# Architecture 4: 3 hidden layers with 4, 8, and 16 neurons
model4 = Sequential()
model4.add(Dense(4, input_dim=30, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model4.add(Dropout(0.4))
model4.add(Dense(8, activation='relu'))
model4.add(Dense(16, activation='relu'))
model4.add(Dense(1, activation='sigmoid'))
```

```

model4.add(Dropout(0.5))
model4.add(Dense(16, activation='relu'))
model4.add(Dense(1, activation='sigmoid'))

# Architecture 5: 3 hidden layers with 8 neurons each
model5 = Sequential()
model5.add(Dense(8, input_dim=30, activation='relu', kernel_regularizer=regularizers.l2(0.01)))
model5.add(Dropout(0.5))
model5.add(Dense(8, activation='relu'))
model5.add(Dropout(0.5))
model5.add(Dense(8, activation='relu'))
model5.add(Dense(1, activation='sigmoid'))

# Define common hyperparameters
epochs = 10
batch_size = 16
verbose = 1
validation_split = 0.2

# Define the learning rates and dropout rates for each architecture
learning_rates = [0.001, 0.002, 0.003, 0.004, 0.005]
dropout_rates = [0.1, 0.2, 0.3, 0.4, 0.5]

# Compile and train each model with the corresponding learning rate and dropout rate
models = [model1, model2, model3, model4, model5]
for i, model in enumerate(models):
    learning_rate = learning_rates[i]
    dropout_rate = dropout_rates[i]
    optimizer = Adam(learning_rate=learning_rate)
    model.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])

    # Define early stopping callback
    early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

    print("Training Architecture", i+1)
    history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=verbose,
                        validation_split=validation_split, callbacks=[early_stopping])

```

```

        validation_split=validation_split, callbacks=[early_stopping])

# Print training and validation accuracy
print("Training accuracy:", history.history['accuracy'][-1])
print("Validation accuracy:", history.history['val_accuracy'][-1])

# Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)

# Make predictions on the test set
y_pred = model.predict(X_test)
y_pred = np.round(y_pred).flatten()

# Compute classification report
report = classification_report(y_test, y_pred)
print(report)

```

Training Architecture 1

Epoch 1/10

22/22 [=====] - 1s 12ms/step - loss: 0.7618 - accuracy: 0.5559 - val\_loss: 0.7457 - val\_accuracy: 0.6163

Epoch 2/10

22/22 [=====] - 0s 5ms/step - loss: 0.7382 - accuracy: 0.6235 - val\_loss: 0.7263 - val\_accuracy: 0.6628

Epoch 3/10

22/22 [=====] - 0s 6ms/step - loss: 0.7243 - accuracy: 0.6794 - val\_loss: 0.7074 - val\_accuracy: 0.7558

Epoch 4/10

22/22 [=====] - 0s 7ms/step - loss: 0.7085 - accuracy: 0.7265 - val\_loss: 0.6897 - val\_accuracy: 0.7791

Epoch 5/10

22/22 [=====] - 0s 7ms/step - loss: 0.6865 - accuracy: 0.7794 - val\_loss: 0.6725 - val\_accuracy: 0.8372

Epoch 6/10

22/22 [=====] - 0s 7ms/step - loss: 0.6732 - accuracy: 0.8059 - val\_loss: 0.6564 - val\_accuracy: 0.8605

```
04 - val_accuracy. 0.8000
Epoch 7/10
```

Confusion matrix of varied architectures

## PERFORMANCE METRICS

```
In [65]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from tabulate import tabulate

# Define a function to plot the confusion matrix
def plot_confusion_matrix(ax, y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)
    ax.set_title(title)
    ax.set_xlabel('Predicted Label')
    ax.set_ylabel('True Label')

# Train and evaluate each model
models = [model1, model2, model3, model4, model5]
accuracy_scores = []
precision_scores = []
recall_scores = []
f1_scores = []

fig, axs = plt.subplots(1, len(models), figsize=(20, 6))

for i, model in enumerate(models):
    print("Training Architecture", i+1)
    history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=verbose,
                        validation_split=validation_split, callbacks=[early_stopping])

    # Print training and validation accuracy
```

```
print("Training accuracy:", history.history['accuracy'][-1])
print("Validation accuracy:", history.history['val_accuracy'][-1])

# Evaluate the model
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)

# Make predictions on the test set
y_pred_prob = model.predict(X_test)
y_pred = np.round(y_pred_prob).flatten()

# Calculate accuracy, precision, recall, and F1 score
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Append scores to lists
accuracy_scores.append(accuracy)
precision_scores.append(precision)
recall_scores.append(recall)
f1_scores.append(f1)

# Plot the confusion matrix
plot_confusion_matrix(axes[i], y_test, y_pred, title=f"Architecture {i+1}")

# Adjust the spacing between subplots
plt.tight_layout()

# Combine the scores and architectures into a table
table = zip(['Architecture 1', 'Architecture 2', 'Architecture 3', 'Architecture 4', 'Architecture 5'],
            accuracy_scores, precision_scores, recall_scores, f1_scores)

# Define the headers for the table
headers = ['Architecture', 'Accuracy', 'Precision', 'Recall', 'F1-Score']
```



```

# Print the table using tabulate
print(tabulate(table, headers=headers, tablefmt='grid'))

# Plot the performance metrics comparison
architectures = ['Architecture 1', 'Architecture 2', 'Architecture 3', 'Architecture 4', 'Architecture 5']
x = np.arange(len(architectures))
width = 0.2

plt.figure(figsize=(10, 8))
plt.bar(x - 1.5 * width, accuracy_scores, width, label='Accuracy')
plt.bar(x - 0.5 * width, precision_scores, width, label='Precision')
plt.bar(x + 0.5 * width, recall_scores, width, label='Recall')
plt.bar(x + 1.5 * width, f1_scores, width, label='F1-Score')

plt.xlabel('Architecture')
plt.ylabel('Score')
plt.title('Performance Metrics Comparison')
plt.xticks(x, architectures)

# Adjust the position of the legend
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')

plt.show()

```

Training Architecture 1

Epoch 1/10

22/22 [=====] - 0s 11ms/step - loss: 0.6067 - accuracy: 0.8529 - val\_loss: 0.5870 - val\_accuracy: 0.8605

Epoch 2/10

22/22 [=====] - 0s 7ms/step - loss: 0.5876 - accuracy: 0.8794 - val\_loss: 0.5751 - val\_accuracy: 0.8721

Epoch 3/10

22/22 [=====] - 0s 13ms/step - loss: 0.5875 - accuracy: 0.8706 - val\_loss: 0.5629 - val\_accuracy: 0.8721

Epoch 4/10

```

22/22 [=====] - 0s 9ms/step - loss: 0.5687 - accuracy: 0.8794 - val_loss: 0.55
17 - val_accuracy: 0.8721
Epoch 5/10
22/22 [=====] - 0s 7ms/step - loss: 0.5593 - accuracy: 0.8824 - val_loss: 0.54
03 - val_accuracy: 0.9070
Epoch 6/10
22/22 [=====] - 0s 6ms/step - loss: 0.5443 - accuracy: 0.8824 - val_loss: 0.52
97 - val_accuracy: 0.9070
Epoch 7/10
22/22 [=====] - 0s 8ms/step - loss: 0.5336 - accuracy: 0.9059 - val_loss: 0.51
92 - val_accuracy: 0.8953
Epoch 8/10
22/22 [=====] - 0s 8ms/step - loss: 0.5338 - accuracy: 0.8676 - val_loss: 0.50
93 - val_accuracy: 0.8953
Epoch 9/10
22/22 [=====] - 0s 6ms/step - loss: 0.5163 - accuracy: 0.8794 - val_loss: 0.49
93 - val_accuracy: 0.8953
Epoch 10/10
22/22 [=====] - 0s 7ms/step - loss: 0.5114 - accuracy: 0.8765 - val_loss: 0.48
95 - val_accuracy: 0.8837
Training accuracy: 0.8764705657958984
Validation accuracy: 0.8837209343910217
5/5 [=====] - 0s 3ms/step - loss: 0.4722 - accuracy: 0.9580
Test accuracy: 0.9580419659614563
5/5 [=====] - 0s 3ms/step
Training Architecture 2
Epoch 1/10
22/22 [=====] - 0s 18ms/step - loss: 0.4579 - accuracy: 0.8441 - val_loss: 0.3
926 - val_accuracy: 0.9186
Epoch 2/10
22/22 [=====] - 0s 13ms/step - loss: 0.4396 - accuracy: 0.8500 - val_loss: 0.3
785 - val_accuracy: 0.8953
Epoch 3/10
22/22 [=====] - 0s 17ms/step - loss: 0.4316 - accuracy: 0.8382 - val_loss: 0.3
529 - val_accuracy: 0.9419

```

```

Epoch 4/10
22/22 [=====] - 0s 12ms/step - loss: 0.4195 - accuracy: 0.8500 - val_loss: 0.3
389 - val_accuracy: 0.9186
Epoch 5/10
22/22 [=====] - 0s 14ms/step - loss: 0.4058 - accuracy: 0.8471 - val_loss: 0.3
258 - val_accuracy: 0.9070
Epoch 6/10
22/22 [=====] - 0s 13ms/step - loss: 0.3820 - accuracy: 0.8588 - val_loss: 0.3
049 - val_accuracy: 0.9535
Epoch 7/10
22/22 [=====] - 0s 13ms/step - loss: 0.3890 - accuracy: 0.8676 - val_loss: 0.2
979 - val_accuracy: 0.9535
Epoch 8/10
22/22 [=====] - 0s 11ms/step - loss: 0.3998 - accuracy: 0.8529 - val_loss: 0.2
893 - val_accuracy: 0.9419
Epoch 9/10
22/22 [=====] - 0s 9ms/step - loss: 0.3655 - accuracy: 0.8676 - val_loss: 0.28
24 - val_accuracy: 0.9419
Epoch 10/10
22/22 [=====] - 0s 13ms/step - loss: 0.3744 - accuracy: 0.8559 - val_loss: 0.2
740 - val_accuracy: 0.9419
Training accuracy: 0.8558823466300964
Validation accuracy: 0.9418604373931885
5/5 [=====] - 0s 3ms/step - loss: 0.2464 - accuracy: 0.9720
Test accuracy: 0.9720279574394226
5/5 [=====] - 0s 10ms/step
Training Architecture 3
Epoch 1/10
22/22 [=====] - 0s 12ms/step - loss: 0.3083 - accuracy: 0.9235 - val_loss: 0.2
441 - val_accuracy: 0.9767
Epoch 2/10
22/22 [=====] - 0s 12ms/step - loss: 0.3064 - accuracy: 0.9059 - val_loss: 0.2
382 - val_accuracy: 0.9419
Epoch 3/10
22/22 [=====] - 0s 11ms/step - loss: 0.2562 - accuracy: 0.9412 - val_loss: 0.1
076 - val_accuracy: 0.9884

```

```

9/0 - val_accuracy: 0.9884
Epoch 4/10
22/22 [=====] - 0s 15ms/step - loss: 0.2360 - accuracy: 0.9382 - val_loss: 0.1
819 - val_accuracy: 0.9884
Epoch 5/10
22/22 [=====] - 0s 13ms/step - loss: 0.2590 - accuracy: 0.9235 - val_loss: 0.1
718 - val_accuracy: 0.9767
Epoch 6/10
22/22 [=====] - 0s 10ms/step - loss: 0.2092 - accuracy: 0.9559 - val_loss: 0.1
627 - val_accuracy: 0.9651
Epoch 7/10
22/22 [=====] - 0s 16ms/step - loss: 0.2034 - accuracy: 0.9471 - val_loss: 0.1
554 - val_accuracy: 0.9884
Epoch 8/10
22/22 [=====] - 0s 17ms/step - loss: 0.2120 - accuracy: 0.9412 - val_loss: 0.1
485 - val_accuracy: 0.9884

Epoch 9/10
22/22 [=====] - 0s 9ms/step - loss: 0.2435 - accuracy: 0.9088 - val_loss: 0.14
69 - val_accuracy: 0.9884
Epoch 10/10
22/22 [=====] - 0s 11ms/step - loss: 0.2309 - accuracy: 0.9206 - val_loss: 0.1
438 - val_accuracy: 0.9651
Training accuracy: 0.9205882549285889
Validation accuracy: 0.9651162624359131
5/5 [=====] - 0s 13ms/step - loss: 0.1232 - accuracy: 0.9720
Test accuracy: 0.9720279574394226
5/5 [=====] - 0s 3ms/step
Training Architecture 4
Epoch 1/10
22/22 [=====] - 0s 13ms/step - loss: 0.4741 - accuracy: 0.8088 - val_loss: 0.4
639 - val_accuracy: 0.8140
Epoch 2/10
22/22 [=====] - 0s 12ms/step - loss: 0.4418 - accuracy: 0.8382 - val_loss: 0.4
641 - val_accuracy: 0.8023
Epoch 3/10
22/22 [=====] - 0s 12ms/step - loss: 0.4050 - accuracy: 0.8850 - val_loss: 0.44

```

```

22/22 [=====] - 0s 6ms/step - loss: 0.4858 - accuracy: 0.8059 - val_loss: 0.44
70 - val_accuracy: 0.8256
Epoch 4/10
22/22 [=====] - 0s 8ms/step - loss: 0.4365 - accuracy: 0.8294 - val_loss: 0.37
73 - val_accuracy: 0.8721
Epoch 5/10
22/22 [=====] - 0s 18ms/step - loss: 0.4929 - accuracy: 0.8000 - val_loss: 0.4
383 - val_accuracy: 0.8256
Epoch 6/10
22/22 [=====] - 0s 9ms/step - loss: 0.4553 - accuracy: 0.8206 - val_loss: 0.40
25 - val_accuracy: 0.8488
Epoch 7/10
22/22 [=====] - 0s 12ms/step - loss: 0.4628 - accuracy: 0.8176 - val_loss: 0.3
689 - val_accuracy: 0.8721
Epoch 8/10
22/22 [=====] - 0s 13ms/step - loss: 0.4407 - accuracy: 0.8324 - val_loss: 0.4
389 - val_accuracy: 0.8256
Epoch 9/10
22/22 [=====] - 0s 14ms/step - loss: 0.4840 - accuracy: 0.7971 - val_loss: 0.4
675 - val_accuracy: 0.8140
Epoch 10/10
22/22 [=====] - 0s 10ms/step - loss: 0.4392 - accuracy: 0.8265 - val_loss: 0.4
204 - val_accuracy: 0.8256
Training accuracy: 0.8264706134796143
Validation accuracy: 0.8255813717842102
5/5 [=====] - 0s 9ms/step - loss: 0.3030 - accuracy: 0.9371
Test accuracy: 0.9370629191398621
5/5 [=====] - 0s 4ms/step
Training Architecture 5
Epoch 1/10
22/22 [=====] - 0s 14ms/step - loss: 0.3136 - accuracy: 0.9000 - val_loss: 0.1
751 - val_accuracy: 0.9651
Epoch 2/10
22/22 [=====] - 0s 16ms/step - loss: 0.3409 - accuracy: 0.8912 - val_loss: 0.2
082 - val_accuracy: 0.9302

```



```


Epoch 3/10
22/22 [=====] - 0s 13ms/step - loss: 0.2787 - accuracy: 0.9206 - val_loss: 0.1
867 - val_accuracy: 0.9535
Epoch 4/10
22/22 [=====] - 0s 10ms/step - loss: 0.3093 - accuracy: 0.8912 - val_loss: 0.1
591 - val_accuracy: 0.9767
Epoch 5/10
22/22 [=====] - 0s 13ms/step - loss: 0.3298 - accuracy: 0.8971 - val_loss: 0.1
632 - val_accuracy: 0.9767
Epoch 6/10
22/22 [=====] - 0s 12ms/step - loss: 0.2338 - accuracy: 0.9353 - val_loss: 0.1
480 - val_accuracy: 0.9884
Epoch 7/10
22/22 [=====] - 0s 10ms/step - loss: 0.2415 - accuracy: 0.9235 - val_loss: 0.1
607 - val_accuracy: 0.9535
Epoch 8/10
22/22 [=====] - 0s 11ms/step - loss: 0.2626 - accuracy: 0.9235 - val_loss: 0.1
690 - val_accuracy: 0.9302
Epoch 9/10
22/22 [=====] - 0s 13ms/step - loss: 0.2757 - accuracy: 0.9118 - val_loss: 0.1
508 - val_accuracy: 0.9651
Training accuracy: 0.9117646813392639
Validation accuracy: 0.9651162624359131
5/5 [=====] - 0s 10ms/step - loss: 0.1353 - accuracy: 0.9650
Test accuracy: 0.9650349617004395
5/5 [=====] - 0s 5ms/step

```

Architecture	Accuracy	Precision	Recall	F1-Score
Architecture 1	0.958042	1	0.888889	0.941176
Architecture 2	0.972028	1	0.925926	0.961538
Architecture 3	0.972028	0.980769	0.944444	0.962264

Architecture 4	0.937063	1	0.833333	0.909091
Architecture 5	0.965035	0.962264	0.944444	0.953271





```
In [66]: # Define a function to plot the confusion matrix
def plot_confusion_matrix(ax, y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)
    ax.set_title(title)
    ax.set_xlabel('Predicted Label')
    ax.set_ylabel('True Label')

# Train and evaluate each model
models = [model1, model2, model3, model4, model5]
test_scores = []

fig, axs = plt.subplots(1, len(models), figsize=(20, 6))
```



```
for i, model in enumerate(models):
    print("Training Architecture", i+1)
    optimizer = Adam(learning_rate=learning_rates[i])
    model.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])

    history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=verbose,
                        validation_split=validation_split, callbacks=[early_stopping])

    # Print training and validation accuracy
    print("Training accuracy:", history.history['accuracy'][-1])
    print("Validation accuracy:", history.history['val_accuracy'][-1])

    # Evaluate the model
    test_loss, test_acc = model.evaluate(X_test, y_test)
    print('Test accuracy:', test_acc)
    test_scores.append(test_acc)

    # Make predictions on the test set
    y_pred = model.predict(X_test)
    y_pred_classes = np.where(y_pred > 0.5, 1, 0)



    # Plot the confusion matrix
    plot_confusion_matrix(axs[i], y_test, y_pred_classes, title=f"Architecture {i+1}")

    # Print precision, recall, and f1-score
    print(classification_report(y_test, y_pred_classes))

# Adjust the spacing between subplots
plt.tight_layout()

# Plot the testing scores comparison
plt.figure(figsize=(4, 2))
plt.plot(range(1, 6), test_scores, marker='o')
plt.xlabel('Architecture')
plt.ylabel('Test Accuracy')
plt.title('Testing Accuracy Comparison')
```

```
plt.xticks(range(1, 6))  
plt.show()
```



# ROC CURVE FOR THE ARCHITECTURES

In [67]:

```

import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

# Function to plot the ROC curve
def plot_roc_curve(model, X_test, y_test, label):
    y_pred = model.predict(X_test).ravel()
    fpr, tpr, thresholds = roc_curve(y_test, y_pred)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label=f"{label} (AUC = {roc_auc:.2f})")

# Plot the ROC curves for all architectures
plt.figure(figsize=(6, 4))
plt.plot([0, 1], [0, 1], 'k--')

models = [model1, model2, model3, model4, model5]
labels = ["Architecture 1", "Architecture 2", "Architecture 3", "Architecture 4", "Architecture 5"]

for model, label in zip(models, labels):
    plot_roc_curve(model, X_test, y_test, label)


plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for Different Architectures')
plt.legend(loc='lower right')
plt.show()


```

```

5/5 [=====] - 0s 2ms/step
5/5 [=====] - 0s 2ms/step
5/5 [=====] - 0s 2ms/step
5/5 [=====] - 0s 2ms/step
5/5 [=====] - 0s 2ms/step

```





IMPACT OF THE HYPERPARAMETERS TUNING

Increasing the number of hidden layers, neurons, epochs, and batch size can have varying effects and consequences on a neural network's performance, depending on its specific task, dataset, and architecture. Here are some general points to consider:

- More hidden layers can help the network learn complex and abstract representations of input data, but it may overfit and memorize the training data instead of generalizing to new examples. To prevent this the network should be validated on a separate set.
- More neurons per layer can increase the network's capacity, but it can also slow down training, increase overfitting, and lead to vanishing or exploding gradients. The number of neurons should be chosen based on the task's complexity and dataset size.
- Increasing epochs allows the network to see more examples and improve its parameters, but it can also lead to overfitting. Validation loss should be monitored, and early stopping techniques can be used.
- Larger batch sizes can speed up training and stabilize optimization, but they can also require more resources, reduce generalization performance, and increase the risk of getting stuck in local minima. Batch size should be chosen based on available resources and dataset/network characteristics.

## OTHER MACHINE LEARNING ALGORITHMS ON BREAST CANCER DATASET

### SUPPORT VECTOR MACHINE & NAIVE BAYES

Support Vector Machine (SVM) algorithm is used to classify or predict outcomes based on input data. SVM works by finding a straight line (called a hyperplane) that separates the input data into different categories or classes. In binary classification like ours, SVM finds the hyperplane that maximizes the distance between the hyperplane and the closest data points from each category.

Define input and target variable


```
In [68]: #dropping column "id", NaN values, and target "label"  
X = df.drop(labels=["label", "id", "Unnamed: 32"], axis=1)
```

```
In [69]: Y = labelencoder.fit_transform(df["label"].values)  
print("Label after encoding are: ", np.unique(Y))
```

Label after encoding are: [0 1]

```
In [70]: sns.scatterplot(x='area_mean',y='smoothness_mean',hue='label',data=df)
plt.ioff()
```

```
Out[70]: <matplotlib.pyplot._IoffContext at 0x7ff14f8e7490>
```



In [71]:

```
#splitting dataset to train and test set  
X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=42)
```

In [72]:

```
#check shape of train  
X_train.shape
```

Out[72]: (455, 30)

In [73]:

```
#check shape of test  
X_test.shape
```

Out[73]: (114, 30)

## SUPPORT VECTOR MACHINE MODEL BUILDING

In [74]:

```
from sklearn.svm import SVC  
from sklearn.metrics import classification_report,confusion_matrix, roc_auc_score, roc_curve  
svc_model=SVC()  
svc_model = SVC(probability=True)  
# Fit the SVM model on the training data  
svc_model.fit(X_train,y_train)
```

Out[74]: SVC(probability=True)

## MODEL EVALUATION



In [75]:

```

# Generate predictions on the test set
y_predict = svc_model.predict(X_test)

# Compute the AUC score
auc = roc_auc_score(y_test, y_predict)
print("AUC Score: {:.2f}".format(auc))

# Print classification report
print(classification_report(y_test, y_predict))

```

AUC Score: 0.93


	precision	recall	f1-score	support
0	0.92	1.00	0.96	71
1	1.00	0.86	0.92	43
accuracy			0.95	114
macro avg	0.96	0.93	0.94	114
weighted avg	0.95	0.95	0.95	114

```
In [76]: # Assuming y_test and y_predict are the true and predicted labels respectively
cm = confusion_matrix(y_test, y_predict)

# Plot the confusion matrix
sns.heatmap(cm, annot=True)

# Add axis labels and title

plt.show()
```



## MODEL OPTIMIZATION

```
In [77]: #find best hyper parameters
from sklearn.model_selection import GridSearchCV
param_grid = {'C':[0.1,1,10,100,1000], 'gamma':[1,0.1,0.01,0.001,0.001], 'kernel':['rbf']}
grid = GridSearchCV(SVC(),param_grid,verbose = 4)
grid.fit(X_train,y_train)
grid.best_params_
grid.best_estimator_
grid_predictions = grid.predict(X_test)
cmG = confusion_matrix(y_test,grid_predictions)
sns.heatmap(cmG, annot=True)
print(classification_report(y_test,grid_predictions))
```

Fitting 5 folds for each of 25 candidates, totalling 125 fits

```
[CV 1/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.637 total time= 0.1s
[CV 2/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 3/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 4/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 5/5] END .....C=0.1, gamma=1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 1/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.637 total time= 0.0s
[CV 2/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 3/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 4/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 5/5] END .....C=0.1, gamma=0.1, kernel=rbf;; score=0.626 total time= 0.0s
[CV 1/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.637 total time= 0.0s
[CV 2/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.626 total time= 0.0s
[CV 3/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.626 total time= 0.0s
[CV 4/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.626 total time= 0.0s
[CV 5/5] END .....C=0.1, gamma=0.01, kernel=rbf;; score=0.626 total time= 0.0s
[CV 1/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.637 total time= 0.0s
[CV 2/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.626 total time= 0.0s
[CV 3/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.626 total time= 0.0s
[CV 4/5] END ....C=0.1, gamma=0.001, kernel=rbf;; score=0.626 total time= 0.0s
```

## HYPERPARAMETER OPTIMIZATION

The last model improvement did not yield the percentage of accuracy. Hence, I created a machine learning pipeline that trains a Support Vector Machine (SVM) classifier using a linear kernel and  $C=1$  hyperparameter. This is to classify new data as either one of two categories (binary classification). The dataset is preprocessed using the StandardScaler function, which standardizes the features by removing the mean and scaling to unit variance to ensure that all features have the same impact on the SVM model.

The SVM classifier used is LinearSVC, which finds the best hyperplane to separate the two classes in a high-dimensional space. The hyperplane is chosen to maximize the distance between the two classes, and the classifier is initialized with a  $C=1$  hyperparameter to control the trade-off between maximizing the margin and minimizing the classification error.

In [78]: *#Building a pipeline using a linear classifier*

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.svm import LinearSVC

svm_clf = Pipeline([
    ("scaler", StandardScaler()),
    ("linear_svc", LinearSVC(C=1, loss="hinge", random_state=42))
])

svm_clf.fit(X_train, y_train)
```

Out[78]: Pipeline(steps=[('scaler', StandardScaler()),  
 ('linear\_svc', LinearSVC(C=1, loss='hinge', random\_state=42))])

```
In [79]: #prediction on few values from training set:
predictions = svm_clf.predict(X_train.iloc[:5])
actual = y_train[:5]

print("Predictions\t", "Actual\t\t")
for index in range(len(predictions)):
    print(predictions[index], "\t\t", actual[index])
```

Predictions	Actual
0	0
1	1
0	0
0	0
0	0

```
In [80]: #building model using linear classifier
from sklearn.svm import SVC

# hyperparameter C
C = 5
alpha = 1 / (C * len(X))

svm_clf = SVC(kernel="linear", C=C)
```

```
In [81]: #standardizing input data
import numpy as np

scaler = StandardScaler()

# pre-process the train and test data
X_train_scaled = scaler.fit_transform(X_train.astype(np.float32))
X_test_scaled = scaler.transform(X_test.astype(np.float32))
```

```
In [82]: # train the model
svm_clf.fit(X_train_scaled, y_train)
```

```
Out[82]: SVC(C=5, kernel='linear')
```

```
In [83]: print("SVC: ", svm_clf.intercept_, svm_clf.coef_)

SVC: [0.11182825] [[-0.14169589 -0.01603748 -0.41509059 -0.26099686 -0.1446300
9 -1.12303578
 0.89064088 2.40044064 -0.27473911 0.43838811 1.90485622 -0.28903005
-0.73747757 1.12785897 0.61471225 0.12299309 -0.94256721 0.53829543
-0.77244493 -1.30536627 1.68589978 1.35857697 0.05551069 1.57397528
-0.05466699 -0.59138542 1.66515095 -0.04341149 1.34119342 0.45562077]]
```

### MODEL EVALUATION ON TRAINING DATA

```
In [84]: # function to print out classification model report
def classification_report(model_name, test, pred, label):
    from sklearn.metrics import precision_score, recall_score
    from sklearn.metrics import accuracy_score, f1_score, roc_auc_score

    print(model_name, ":\n")

    print("Accuracy Score: ", '{:,.3f}'.format(float(accuracy_score(test, pred)) * 100), "%")
    print("Precision: ", '{:,.3f}'.format(float(precision_score(test, pred, pos_label=label)) * 100), "%")
    print("Recall: ", '{:,.3f}'.format(float(recall_score(test, pred, pos_label=label)) * 100), "%")
    print("F1 score: ", '{:,.3f}'.format(float(f1_score(test, pred, pos_label=label)) * 100), "%")
    print("AUC-ROC: ", '{:,.3f}'.format(float(roc_auc_score(test, pred)) * 100), "%")
```

```
In [85]: svm_clf_pred = svm_clf.predict(X_train_scaled)
classification_report("SVM with linear kernel and C=5 Hyperparameter", y_train, svm_clf_pred, 0)
```

SVM with linear kernel and C=5 Hyperparameter :

```
Accuracy Score: 98.901 %
Precision: 98.616 %
Recall: 99.650 %
F1 score: 99.130 %
AUC-ROC: 98.642 %
```

SVM MODEL EVALUATION ON TEST DATA

```
In [86]: svm_clf_pred_test = svm_clf.predict(X_test_scaled)
```

```
In [87]: classification_report("SVM on Test Set", y_test, svm_clf_pred_test, 0)
```

SVM on Test Set :

```
Accuracy Score: 96.491 %
Precision: 98.551 %
Recall: 95.775 %
F1 score: 97.143 %
AUC-ROC: 96.725 %
```

NAIVE BAYES ALGORITHM ON BREAST CANCER DATASET

Define input and target variable



```
In [88]: #dropping column "id", NaN values, and target "label"
X = df.drop(labels=["label", "id", "Unnamed: 32"], axis=1)
Y = labelencoder.fit_transform(df["label"].values)
print("Label after encoding are: ", np.unique(Y))
```

Label after encoding are: [0 1]

Dataset is normalized to adjust the scale of input features, so that they are all on a similar range. This helps to avoid situations where certain features have a larger impact on the model than others, which can cause the model to make biased predictions.

```
In [89]: # Normalization of dataset
X = (X - np.min(X)) / (np.max(X) - np.min(X))
```

```
/opt/anaconda3/lib/python3.9/site-packages/numpy/core/fromnumeric.py:84: FutureWarning: In a future ver
sion, DataFrame.min(axis=None) will return a scalar min over the entire DataFrame. To retain the old be
havior, use 'frame.min(axis=0)' or just 'frame.min()'
  return reduction(axis=axis, out=out, **passkwargs)
/opt/anaconda3/lib/python3.9/site-packages/numpy/core/fromnumeric.py:84: FutureWarning: In a future ver
sion, DataFrame.max(axis=None) will return a scalar max over the entire DataFrame. To retain the old be
havior, use 'frame.max(axis=0)' or just 'frame.max()'
  return reduction(axis=axis, out=out, **passkwargs)
```

```
In [90]: #Splitting data to train and test set
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.3, random_state = 42)
```

Bernoulli Naive Bayes

The Bernoulli Naive Bayes (BNB) algorithm is used to classify input data that consists of binary or boolean features, where each feature can take on one of two possible values: 0 or 1. The algorithm assumes that each feature is independent of all the other features given the class variable.

In [91]:

```

111 1911: from sklearn.metrics import classification_report, roc_auc_score
from sklearn.naive_bayes import BernoulliNB

BNB = BernoulliNB()
BNB.fit(X_train, y_train)
y_pred = BNB.predict(X_test)
y_prob = BNB.predict_proba(X_test)[: , 1]

print("Accuracy:", BNB.score(X_test, y_test))
print(classification_report(y_test, y_pred, labels=[0, 1]))
print("AUC Score:", roc_auc_score(y_test, y_prob))

```

Accuracy: 0.631578947368421

	precision	recall	f1-score	support
0	0.63	1.00	0.77	108
1	0.00	0.00	0.00	63
accuracy			0.63	171
macro avg	0.32	0.50	0.39	171
weighted avg	0.40	0.63	0.49	171

AUC Score: 0.5324074074074073

/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

## MODEL OPTIMIZATION OF THE NAIVE BAYES MODEL

## Multinomial Naive Bayes

Multinomial Naive Bayes (MNB) algorithm is used to classify input data consisting of count or discrete data such as text classification. This means that the input features represent the count of a certain event or occurrence.

```
In [92]: #Building the Multinomial Naive Bayes Classifier Model
from sklearn.naive_bayes import MultinomialNB
MNB = MultinomialNB()
MNB.fit(X_train, y_train)
MNB.score(X_test, y_test)
y_pred = MNB.predict(X_test)
y_prob = MNB.predict_proba(X_test)[: , 1]
print("Accuracy:", MNB.score(X_test, y_test))
print(classification_report(y_test, y_pred))
print("AUC Score:", roc_auc_score(y_test, y_prob))
```

Accuracy: 0.8304093567251462

	precision	recall	f1-score	support
0	0.79	1.00	0.88	108
1	1.00	0.54	0.70	63
accuracy			0.83	171
macro avg	0.89	0.77	0.79	171
weighted avg	0.87	0.83	0.82	171

AUC Score: 0.9528218694885362

## GUASSIAN NAIVE BAYES

Gaussian Naive Bayes (GNB) algorithm is used to classify input data where the continuous-valued features of each class are assumed to be normally distributed. GNB assumes that the distribution of the features is normal. GNB is commonly used when dealing with continuous data and when the distribution of the features is assumed to be Gaussian.

```
In [93]: #Building the Naive Bayes Classifier Model
from sklearn.naive_bayes import GaussianNB
GNB = GaussianNB()
GNB.fit(X_train, y_train)
print("Naive Bayes score: ", GNB.score(X_test, y_test))
y_pred = GNB.predict(X_test)
y_prob = GNB.predict_proba(X_test)[:, 1]
print("Accuracy:", GNB.score(X_test, y_test))
print(classification_report(y_test, y_pred))
print("AUC Score:", roc_auc_score(y_test, y_prob))
```

Naive Bayes score: 0.935672514619883

Accuracy: 0.935672514619883

	precision	recall	f1-score	support
0	0.94	0.95	0.95	108
1	0.92	0.90	0.91	63
accuracy			0.94	171
macro avg	0.93	0.93	0.93	171
weighted avg	0.94	0.94	0.94	171

AUC Score: 0.9926513815402704

## REFERENCE

kaggle.com. (n.d.). Naive Bayes Implementation on Cancer Dataset. [online] Available at: <https://www.kaggle.com/code/nisasoylu/naive-bayes-implementation-on-cancer-dataset> (<https://www.kaggle.com/code/nisasoylu/naive-bayes-implementation-on-cancer-dataset>) [Accessed 14 May 2023].

Eren, M.E. (2020). Support Vector Machines on the Breast Cancer Wisconsin (Diagnostic) Data Set. [online] GitHub. Available at: [https://github.com/MaksimEkin/Breast-Cancer-Prediction-SVM/blob/master/breast\\_cancer\\_prediction.ipynb](https://github.com/MaksimEkin/Breast-Cancer-Prediction-SVM/blob/master/breast_cancer_prediction.ipynb) ([https://github.com/MaksimEkin/Breast-Cancer-Prediction-SVM/blob/master/breast\\_cancer\\_prediction.ipynb](https://github.com/MaksimEkin/Breast-Cancer-Prediction-SVM/blob/master/breast_cancer_prediction.ipynb)) [Accessed 14 May 2023].